

APH101-Biostatistics And R

Yu.Lu23

2025-06-02

Table of contents

Final review of statistics knowledge	4
Three main sampling distributions in hypothesis testing	4
Chi-square distribution	4
Properties	5
Application	5
Student's t-distribution	5
Construction	5
Properties	6
F-distribution	6
Example	6
Estimation of Population Characteristic	6
Point Estimation	6
Interval Estimation—Confidence interval(CI)	6
Definition	7
Interpretation	7
CI on Mean	7
CI for Estimating σ	8
Hypothesis Testing	9
Example of one-sample t-test	10
a	10
b	10
c	11
d	11
Example of unpaired t-test	12
Example of one-sample Variance Test	12
Example of two-sample Variance Test	13

ANOVA– Analysis of Variance	13
Example of two-way ANOVA	15
Non-parametric tests	20
Application of testing the goodness of fit	20
Application of testing for homogeneity	22
Application of testing for independence	24
Linear Regression	27
Simple linear regression	27
Brief introduction	27
Maximum Likelihood Estimation (MLE)	28
Hypothesis testing for estimates with unknown σ^2	29
R^2 –the fraction of variability explained by the regression	32
Multiple linear regression (MLR)	32
A potential problem in practice –multicollinearity	33
Adjusted R-squared	33
Logistic Regression	34
Odds	36
Generalized Linear Models	36
Random Component	37
Exponential Family	37
Systematic Component and Link Component	37
Example	38
Gaussian-noise Linear Regression	38
Bernoulli	38
Link Function	38
Logistic	39
Survival Analysis	39
Some definations	39
Kaplan-Meier estimate	40
Cox-proportional hazards model	40
Final review of R codes	42
Calculation	42
Vectors	42
Rmd knowledge	45
Probability in R	45
R basic	48

Data class	48
Numeric	48
Integer	48
Character	49
Logical	49
Factor	50
Date and time	51
Data frame	52
List	53
Matrix	55
Example 1	57
Data visualization (mainly: ggplot2)	61
Box plots	61
Histogram plots	63
Density plots	66
Scatter plots	67
Scatter plots with regression line	68
Heat map	71
Faceting (make many panels of graphics where each panel represents the same relationship between variables, but something changes between each pane)	74
Data manipulation	80
apply	81
Tibbles	82
%>%	83
select	83
filter	83
slice	84
arrange	84
mutate	85
summarise	86
group_by	86
examples	87
ps:the comparison between whether to use %>% or not	92
Control flow	93
while loop	93
for loop	93
Fibonacci sequence	93
bootstrap estimate of a sampling distribution	94

Functions	95
Functions construction	95
Ellipses	95
Linear regression and multiple linear regression (Lab 10)	96
Logistic regression	115
Simple example	115
Confusion Matrix	116
example	116
Lung Cancer Classification (https://www.kaggle.com/datasets/mysarahmadbhat/lung-cancer?select=survey+lung+cancer.csv)	116

Final review of statistics knowledge

Three main sampling distributions in hypothesis testing

Chi-square distribution

Suppose $X_1, \dots, X_n \stackrel{\text{i.i.d.}}{\sim} \mathcal{N}(0, 1)$. the distribution of the statistic

$$X_1^2 + \dots + X_n^2$$

is called a chi-square distribution with n degrees of freedom, denoted by $\chi^2(n)$.

Besides, random variable $X_i^2 \sim \text{Gamma}(\frac{1}{2}, \frac{1}{2})$ corresponds to the chi-squared distribution with 1 degree of freedom, denoted as χ_1^2 .

This is derived by the MGF:

Since

$$M_{X_1^2 + \dots + X_n^2}(t) = M_{X_1^2}(t) \times \dots \times M_{X_n^2}(t) = \begin{cases} \infty & t \geq \frac{1}{2} \\ (1 - 2t)^{-\frac{n}{2}} & t < \frac{1}{2} \end{cases}$$

This is the MGF of the Gamma $(\frac{n}{2}, \frac{1}{2})$ distribution, so $X_1^2 + \dots + X_n^2 \sim \text{Gamma}(\frac{n}{2}, \frac{1}{2})$. This is called the chi-squared distribution with n degree of freedom, denoted χ_n^2 .

Properties

- If W_1, \dots, W_n are independent χ^2 random variables with, respectively, v_1, \dots, v_n degrees of freedom, then the random variable $W_1 + \dots + W_n$ follows a χ^2 -distribution with $v_1 + \dots + v_n$ degree of freedom.
- The random variable $\frac{(\bar{X} - \mu)^2}{\sigma^2/n}$ follows a χ^2 -distribution with 1 degree of freedom when X follows a normal distribution with mean μ and variance σ^2 .

Application

Chi-square distribution is primarily used in testing:

- Goodness-of-fit
- Independence in contingency tables

Student's t-distribution

Construction

The statistic $T = \frac{\bar{X} - \mu}{S/\sqrt{n}}$ follows a t -distribution with $v = n - 1$ degrees of freedom when X_1, \dots, X_n are i.i.d. normal RVs.

$$\bar{X} \sim \mathcal{N}\left(\mu, \frac{\sigma^2}{n}\right) \quad \frac{\bar{X} - \mu}{\sigma/\sqrt{n}} \sim \mathcal{N}(0, 1) \quad \frac{(n-1)s^2}{\sigma^2} \sim \chi_{n-1}^2$$

If we know the population variance σ^2 , we can easily do inference using the statistic $\frac{\bar{X} - \mu}{\sigma/\sqrt{n}}$. However, σ^2 is usually unknown in practice.

$$\bar{X} \sim \mathcal{N}\left(\mu, \frac{\sigma^2}{n}\right) \quad \frac{\bar{X} - \mu}{\sigma/\sqrt{n}} \sim \mathcal{N}(0, 1) \quad \frac{(n-1)s^2}{\sigma^2} \sim \chi_{n-1}^2$$

We can construct the t -statistic using the sample variance S^2 :

$$T = \frac{\frac{\bar{X} - \mu}{\sigma/\sqrt{n}}}{\sqrt{\frac{(n-1)s^2}{\sigma^2}/(n-1)}} = \frac{\bar{X} - \mu}{S/\sqrt{n}}$$

Notice the sample mean \bar{X} and the sample variance S^2 are independent (the proof is beyond the scope of this course). So the T is now a ratio of a standard normal variable and the square root of a χ^2 RV divided by its degrees of freedom. This is the definition of a t -distribution with $n - 1$ degrees of freedom.

Properties

The t -distribution is primarily used in contexts where the underlying population is assumed to be normally distributed, especially when the sample size is small. Used extensively in problems that deal with inference about population mean μ when population variance σ^2 is unknown; problems where one is trying to determine if means from two samples are significantly different when population variances σ_1^2 and σ_2^2 are unknown.

F-distribution

Let U and V be two independent random variables following χ^2 distributions with ν_1 and ν_2 degrees of freedom, respectively. Then the distribution of the random variable $F = \frac{U/\nu_1}{V/\nu_2}$ is known as F -distribution.

Example

If S_1^2 and S_2^2 are the variances of independent RVs of size n_1 and n_2 taken from normal populations with variances σ_1^2 and σ_2^2 respectively, then

$$F = \frac{S_1^2/\sigma_1^2}{S_2^2/\sigma_2^2} = \frac{\sigma_2^2 S_1^2}{\sigma_1^2 S_2^2}$$

follows an F -distribution with $\nu_1 = n_1 - 1$ and $\nu_2 = n_2 - 1$ degrees of freedom.

Estimation of Population Characteristic

Point Estimation

A point estimate of a population characteristic is a single number that is based on sample data and represents a plausible value of the characteristic.

Interval Estimation—Confidence interval(CI)

An interval estimate of a parameter θ is an interval of the form $\hat{\theta}_L < \theta < \hat{\theta}_U$, where $\hat{\theta}_L$ and $\hat{\theta}_U$ depend on the value of $\hat{\theta}$ for a particular sample and also on the sampling distribution of $\hat{\Theta}$.

- If we were to construct a 95% confidence interval for some population characteristics (population proportion p or population mean μ), we would be using a method that is successful 95% of the time.
- This is also about the question relevant to “How to choose a sample size”

Definition

A $100(1 - \alpha)\%$ confidence interval is an interval of the form $\hat{\theta}_L < \theta < \hat{\theta}_U$, where $\hat{\theta}_L$ and $\hat{\theta}_U$ are respectively values of $\hat{\Theta}_L$ and $\hat{\Theta}_U$ obtained for a particular sample, based on

$$P(\hat{\Theta}_L < \theta < \hat{\Theta}_U) = 1 - \alpha \quad ; \quad 0 < \alpha < 1$$

in the estimation of population parameter θ .

Interpretation

For confidence level of 95% CI for any normal distribution: About 95% of the values are within 1.96 standard deviations of the mean. (Recall the concept of Z-scores)

That is, if

$$\text{Estimate} \pm (Z \times \sigma)$$

was used to generate an interval estimate over and over again with different samples, in the long run 95% of the resulting intervals would include the actual value of the characteristic being estimated.

The confidence level 95% refers to the method used to construct the interval rather than to any particular interval, such as the one we obtained.

CI on Mean

Here, \bar{x} is the sample mean from a simple random sample.

μ is the population mean which we are interested in estimating.

CI on μ with σ known $n \geq 30$ or the population is normal — Use z-statistics

CI: $\bar{x} \pm z_{\alpha/2} \frac{\sigma}{\sqrt{n}}$, for example, 95% CI is $\bar{x} \pm 1.96 \frac{\sigma}{\sqrt{n}}$

One-side Confidence Bound on μ with σ known $n \geq 30$ or the population is normal — Use z-statistics

Upper one-side bound: $\mu < \bar{x} + z_{\alpha} \frac{\sigma}{\sqrt{n}}$

Lower one-side bound: $\mu > \bar{x} - z_{\alpha} \frac{\sigma}{\sqrt{n}}$

For example, 95% Confidence bound on μ is $\bar{x} \pm 1.645 \frac{\sigma}{\sqrt{n}}$

CI on μ with σ unknown and the population is normal — Use t-statistics (use s as the estimate for σ (t-statistics with df = n-1))

CI: $\bar{x} \pm t_{\alpha/2} \frac{s}{\sqrt{n}}$, for example, 95% CI is $\bar{x} \pm t_{0.025} \frac{s}{\sqrt{n}}$ and df = n-1

Remark: The distribution of t is more spread out than the standard normal distribution but when $n \geq 30$, t and z are very close to each other.

CI for $\mu_1 - \mu_2$, both σ_1^2 and σ_2^2 are known

CI of $\mu_1 - \mu_2$: $(\bar{x}_1 - \bar{x}_2) \pm z_{\alpha/2} \sqrt{\frac{\sigma_1^2}{n_1} + \frac{\sigma_2^2}{n_2}}$

CI for $\mu_1 - \mu_2$, both σ_1^2 and σ_2^2 are unknown but assumed equal

CI of $\mu_1 - \mu_2$: $(\bar{x}_1 - \bar{x}_2) \pm t_{\alpha/2} s_p \sqrt{\frac{1}{n_1} + \frac{1}{n_2}}$ with df = $n_1 + n_2 - 2$ where $s_p = \sqrt{\frac{(n_1-1)s_1^2 + (n_2-1)s_2^2}{n_1 + n_2 - 2}}$

CI for paired observations

Previous, we have two independent samples, now we have two dependent samples. We can use the difference between the two samples to construct a confidence interval.

CI of μ_d : $\bar{d} \pm t_{\alpha/2} \frac{s_d}{\sqrt{n}}$ with df = n-1 where s_d is the sample standard deviation of the differences $d_i = x_{1i} - x_{2i}$ and \bar{d} is the sample mean of the differences.

CI for Estimating σ

a $100(1 - \alpha)\%$ CI for σ^2 is $\left(\frac{(n-1)S^2}{\chi_{\alpha/2, n-1}^2}, \frac{(n-1)S^2}{\chi_{1-\alpha/2, n-1}^2} \right)$

where S^2 is the sample variance and $\chi_{\alpha/2, n-1}^2$ and $\chi_{1-\alpha/2, n-1}^2$ are the critical values of the chi-square distribution with $n - 1$ degrees of freedom.

Estimating σ_1^2/σ_2^2

A $100(1 - \alpha)\%$ CI for $\frac{\sigma_1^2}{\sigma_2^2}$ using F-statistics with $f_{1-\alpha/2}(n_1 - 1, n_2 - 1) = 1/f_{\alpha/2}(n_1 - 1, n_2 - 1)$ is

$$\frac{s_1^2}{s_2^2} \frac{1}{f_{\alpha/2}(n_1 - 1, n_2 - 1)} < \frac{\sigma_1^2}{\sigma_2^2} < \frac{s_1^2}{s_2^2} f_{\alpha/2}(n_2 - 1, n_1 - 1)$$

where $f_{\alpha/2}(v_1, v_2)$ is an F-value with v_1 and v_2 degrees of freedom, leaving an area of $\alpha/2$ to the right, and $f_{\alpha/2}(v_2, v_1)$ is a similar F-value with v_2 and v_1 degrees of freedom.

Hypothesis Testing

- z-test

Suppose $X_1, \dots, X_n \xrightarrow[\sim]{\text{iid}} \mathcal{N}(\mu, \sigma^2)$, where μ is unknown and where $\sigma^2 = \sigma_0^2$ is known.

Suppose we wish to test $H_0 : \mu = \mu_0$ against $H_1 : \mu > \mu_0$. Then we can use the test statistic

$$Z = \frac{\bar{X} - \mu_0}{\sigma_0/\sqrt{n}}$$

If H_0 is true then $Z \sim \mathcal{N}(0, 1)$. Let

$$z_{\text{obs}} = \frac{\bar{x} - \mu_0}{\sigma_0/\sqrt{n}}$$

A large value of z_{obs} casts doubt on the validity of H_0 and indicates a departure from H_0 in the direction of H_1 . So the p -value for testing H_0 against H_1 is

$$\begin{aligned} p &= P(Z \geq Z_{\text{obs}} \mid H_0) \\ &= P(\mathcal{N}(0, 1) \geq Z_{\text{obs}}) \\ &= 1 - \Phi(Z_{\text{obs}}) \end{aligned}$$

The z-test of $H_0 : \mu = \mu_0$ against the alternative $H_1 : \mu < \mu_0$ is similar but this time a small, i.e. very negative, value of Z_{obs} casts doubt on H_0 . So the p -value is

$$\begin{aligned} p &= P(Z \leq Z_{\text{obs}} \mid H_0) \\ &= P(\mathcal{N}(0, 1) \leq Z_{\text{obs}}) \\ &= \Phi(Z_{\text{obs}}) \end{aligned}$$

Finally, consider testing $H_0 : \mu = \mu_0$ against the alternative $H_1 : \mu \neq \mu_0$. Let $z_0 = |z_{\text{obs}}|$. A large value of z_0 indicates a departure from H_0 , so the p -value is

$$\begin{aligned} p &= P(|Z| \geq z_0 \mid H_0) \\ &= P(\mathcal{N}(0, 1) \geq z_0) + P(\mathcal{N}(0, 1) \leq -z_0) \\ &= 2(1 - \phi(z_0)) \end{aligned}$$

- t-test

We can use the test statistic

$$T = \frac{\bar{X} - \mu_0}{S/\sqrt{n}}$$

If H_0 is true then $T \sim t_{n-1}$. Let $t_{\text{obs}} = t(x) = \frac{\bar{x} - \mu_0}{s/\sqrt{n}}$ and $t_0 = |t_{\text{obs}}|$. Similarly, we have \square the p -value is $P(t_{n-1} \geq t_{\text{obs}})$ - the p -value is $P(t_{n-1} \leq t_{\text{obs}})$ the p -value is $2P(t_{n-1} \geq t_0)$

Example of one-sample t-test

A marine biologist is studying a species of fish known to have an average length of 20 cm in ocean populations. A new population in a freshwater lake is being analyzed to determine if the environmental differences have altered the fish's average length. The biologist measures the lengths of 10 randomly selected fish, yielding the following data:

22, 23, 21, 24, 22, 20, 25, 19, 23, 22

Assuming the data satisfy the assumption of normality, please address the following using a significance level of 0.1:

a

- null hypothesis: The mean length of fish is 20 cm ($H_0 : \mu = 20$).
- alternative hypothesis: The mean length of fish is not 20 cm ($H_1 : \mu \neq 20$).

b

Since the data is supposed to be normally distributed, the sampling distribution of the sample mean follows t-distribution. The t-test statistic is calculated as follows:

$$t = \frac{\bar{x} - \mu}{s/\sqrt{n}}$$

The t-test statistic is calculated as follows:

```
data2_4 <- c(22, 23, 21, 24, 22, 20, 25, 19, 23, 22)

x_bar <- mean(data2_4)

s <- sd(data2_4)

t <- (x_bar - 20) / (s / sqrt(length(data2_4)))

t
```

```
[1] 3.705882
```

The t-test statistic is 3.705882...

c

Using `pt()` function, the p-value is calculated as follows:

```
p_value <- 2 * pt(-t, df = 9)
p_value
```

```
[1] 0.004875954
```

The p-value is approximately 0.005. Since the p-value is less than 0.05, we reject the null hypothesis.

Therefore, there is sufficient evidence to conclude that the population mean is not equal to 20 which means the environmental differences have altered the fish's average length.

d

Using `qt()` function to find the critical value for a two-tailed test with 90% confidence level:

```
t_critical <- qt(0.95, df = 9)
t_critical
```

```
[1] 1.833113
```

The critical value for a two-tailed test with 90% confidence level is about 1.833113.

Since the t-test statistic 3.705882 is greater than the critical value 1.833113, which is in the critical region. Therefore, we reject the null hypothesis.

Also, we could use confidence interval to verify the result. The 90% confidence interval for the population mean is calculated as follows:

```
ci_4 <- c(x_bar - t_critical * s / sqrt(10), x_bar + t_critical * s / sqrt(10))
ci_4
```

```
[1] 21.06124 23.13876
```

So the 90% confidence interval for the population mean is about (21.1, 23.2).

The confidence interval does not contain the hypothesized value 20. Therefore, we reject the null hypothesis.

Therefore, there is sufficient evidence to conclude that the population mean is not equal to 20 which means the environmental differences have altered the fish's average length.


Example of unpaired t-test

Suppose σ_1^2 and σ_2^2 are unknown but assumed equal. We want to test the null hypothesis $H_0 : \mu_1 = \mu_2$ against the alternative hypothesis $H_1 : \mu_1 \neq \mu_2$. The test statistic is given by

$$t = \frac{\bar{x}_1 - \bar{x}_2}{s_p \sqrt{\frac{1}{n_1} + \frac{1}{n_2}}}$$

where s_p is the pooled sample standard deviation, given by

$$s_p = \sqrt{\frac{(n_1 - 1)s_1^2 + (n_2 - 1)s_2^2}{n_1 + n_2 - 2}}$$



Example of one-sample Variance Test

χ^2 -test for variance. Suppose X_1, \dots, X_n are i.i.d. normal random variables with mean μ and variance σ^2 . We want to test the null hypothesis $H_0 : \sigma^2 = \sigma_0^2$ against the alternative hypothesis $H_1 : \sigma^2 \neq \sigma_0^2$. The test statistic is given by

$$\chi^2 = \frac{(n - 1)s^2}{\sigma^2}$$

Example of two-sample Variance Test

H_0	Test Statistic	H_1	Rejection Region
$\sigma_1^2 = \sigma_2^2$	$f = \frac{s_1^2}{s_2^2}$	$\sigma_1^2 < \sigma_2^2$	$f < f_\alpha(\nu_1, \nu_2)$
	$\sigma_1^2 \neq \sigma_2^2$	$\sigma_1^2 > \sigma_2^2$	$f > f_{1-\alpha}(\nu_1, \nu_2)$
		$f < f_{\alpha/2}(\nu_1, \nu_2)$ or $f > f_{1-\alpha/2}(\nu_1, \nu_2)$	

$\nu_1 = n_1 - 1$ and
 $\nu_2 = n_2 - 1$ are two
degree of freedom.

$$F = \frac{\frac{(n_1-1)S_1^2}{\sigma_1^2} / (n_1 - 1)}{\frac{(n_2-1)S_2^2}{\sigma_2^2} / (n_2 - 1)} = \frac{\sigma_2^2 S_1^2}{\sigma_1^2 S_2^2}$$

If $\sigma_1^2 = \sigma_2^2$, we have

$$F = \frac{S_1^2}{S_2^2} \sim F_{n_1-1, n_2-1}$$

ANOVA– Analysis of Variance

- one-way ANOVA

we need to test the null hypothesis that the group population means are all the same against the alternative that at least one group population mean differs from the others. That is,

$H_0 : \mu_1 = \mu_2 = \dots = \mu_k$ against $H_1 : \text{at least one } \mu_i \text{ differs from the others.}$

ANOVA Table

Source	DF	Sum Sq	Mean Sq	F value	p value
Factor	m-1	11.84 (SS between)	2.9587 (MSB)	8.074 (MSB/MSW)	5.38e – 05 (p-value)
Error	n-m	16.49 (SS Within)	0.3664 (MSW)		
Total	n-1	28.33 (SS Total)			

Source means “the source of the variation in the data.” the possible sources for a one-factor study are Factor, Residuals, and Total.

Factor means “the variability due to the factor of interest.” In the drug example, the factor was the different drug. In the learning example on the previous page, the factor was the method of learning. Sometimes the row heading is labeled as Between.

Error (or Residuals) means “the variability within the groups” or “unexplained random error.” Sometimes the row heading is labeled as Within.

Total means “the total variation in the data from the grand mean”.

DF means “the degrees of freedom in the source.”

Sum Sq means “the sum of squares due to the source.”

Mean Sq means “the mean sum of squares due to the source.”

F value means “the F-statistic.”

P value means “the P-value.”

$SS(\text{Total}) = SS(\text{Between}) + SS(\text{Within})$, where

$SS(\text{Between})$ is the sum of squares between the group means and the grand mean. As the name suggests, it quantifies the variability between the groups of interest.

$SS(\text{Within})$ is the sum of squares between the data and the group means. It quantifies the variability within the groups of interest.

$SS(\text{Total})$ is the sum of squares between the n data points and the grand mean. As the name suggests, it quantifies the total variability in the observed data.

- two-way ANOVA

We can extend the idea of a one-way ANOVA, which tests the effects of one factor on a response variable, to a two-way ANOVA which tests the effects of two factors and their interaction on a response variable.

Source	DF	Sum Sq	MSW	F
Cells	$ab - 1$	$\sum_{i=1}^a \sum_{j=1}^b n (\bar{X}_{ij} - \bar{X}_{....})^2$		
A	$a - 1$	$\sum_{i=1}^a (\bar{X}_{i..} - \bar{X}_{....})^2$	$\frac{SS(A)}{DF(A)}$	$\frac{MS(A)}{MS(\text{Error})}$
B	$b - 1$	$\sum_{j=1}^b (\bar{X}_{.j.} - \bar{X}_{....})^2$	$\frac{SS(B)}{DF(B)}$	$\frac{MS(B)}{MS(\text{Error})}$
$A \times B$	$(a - 1)(b - 1)$	$SS(\text{Cells}) - SS(A) - SS(B)$	$\frac{SS(AB)}{DF(A \times B)}$	$\frac{MS(AB)}{MS(\text{Error})}$
Error	$ab(n - 1)$	$\sum_{i=1}^a \sum_{j=1}^b \sum_{l=1}^n (X_{ijl} - \bar{X}_{ij.})^2$	$\frac{SS(\text{Error})}{DF(\text{Error})}$	
Total	$abn - 1$	$\sum_{i=1}^a \sum_{j=1}^b n (X_{ijl} - \bar{X}_{....})^2$		

- $F = \frac{MS(A)}{MS(Error)}$, for H_0 : no effect of factor A on response variable,
- $F = \frac{MS(B)}{MS(Error)}$, for H_0 : no effect of factor B on response variable,
- $F = \frac{MS(A \times B)}{MS(Error)}$, for H_0 : no effect of interaction on response variable.

We reject any H_0 if $F \geq F_{critical}$; otherwise, we do not reject H_0 .

Example of two-way ANOVA

Two-way ANOVA. In this question, we will use the built-in R data set ToothGrowth to perform two-way ANOVA test. ToothGrowth includes information from a study on the effects of vitamin C on tooth growth in Guinea pigs. The response is the length of odontoblasts (cells responsible for tooth growth) in 60 guinea pigs. Each animal received one of three dose levels of vitamin C (0.5, 1, and 2 mg/day) by one of two delivery methods, orange juice or ascorbic acid (a form of vitamin C and coded as VC). Assuming the data satisfy the assumptions of normality and equal variance, please address the following using a significance level of 0.05

a

- The effects of vitamin C on tooth growth in guinea pigs:
 - null hypothesis: H_0 : mean tooth growth for all doses of vitamin C are equal
 - alternative hypothesis: H_1 : at least one of the means of all doses of vitamin C is different from the others
- The effects of delivery method on tooth growth in guinea pigs:
 - null hypothesis: H_0 : mean tooth growth for the delivery method of orange juice and ascorbic acid are equal.
 - alternative hypothesis H_1 : mean tooth growth for the delivery method of orange juice and ascorbic acid are different.
- The interaction effects of the dose of vitamin C and delivery method on tooth growth in guinea pigs:
 - null hypothesis: H_0 : there is no interaction between the dose of vitamin C and delivery method on tooth growth in guinea pigs, meaning that the relationship between vitamin C and tooth growth is the same for both delivery methods (similarly, the relationship between delivery method and tooth growth is the same for all doses of vitamin C).

- alternative hypothesis: H_1 : there is an interaction between the dose vitamin C and delivery method on tooth growth in guinea pigs, meaning that the relationship between vitamin C and tooth growth is different for both delivery methods (similarly, the relationship between delivery method and tooth growth depends on the dose of vitamin C).

b

We can plot the relationship one by one using two plots


```
library(ggplot2)
data(ToothGrowth)
head(ToothGrowth)
```

```
      len supp dose
1   4.2   VC  0.5
2  11.5   VC  0.5
3   7.3   VC  0.5
4   5.8   VC  0.5
5   6.4   VC  0.5
6  10.0   VC  0.5
```

```
# potential effects of vitamin C on tooth growth.
```

```
ggplot(ToothGrowth, aes(x = factor(dose), y = len)) +
  geom_boxplot() +
  labs(x = "Dose (mg/day)", y = "Tooth Growth (len)", title = "Tooth Growth by Dose of vitamin C")
```



Tooth Growth by Dose of vitamin C



```
# potential effects of delivery method on tooth growth.
```

```
ggplot(ToothGrowth, aes(x = supp, y = len)) +  
  geom_boxplot() +  
  labs(x = "Delivery Method", y = "Tooth Growth (len)", title = "Tooth Growth by Delivery Method")
```

Tooth Growth by Delivery Method



or just one:

```
library(ggplot2)

# potential effects of vitamin C and delivery method.

# OJ represents orange juice and VC represents ascorbic acid.

ggplot(ToothGrowth, aes(x = factor(dose), y = len, fill = supp)) +
  geom_boxplot() +
  labs(x = "Dose (mg/day)", y = "Tooth Growth (len)", title = "Tooth Growth by Dose and Deliv
```



c

```
# Perform two-way ANOVA

anova_result <- aov(len ~ supp * dose, data = ToothGrowth)

summary(anova_result)
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
supp	1	205.4	205.4	12.317	0.000894 ***

```
dose          1 2224.3  2224.3 133.415 < 2e-16 ***
supp:dose     1   88.9    88.9   5.333 0.024631 *
Residuals    56  933.6    16.7
```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

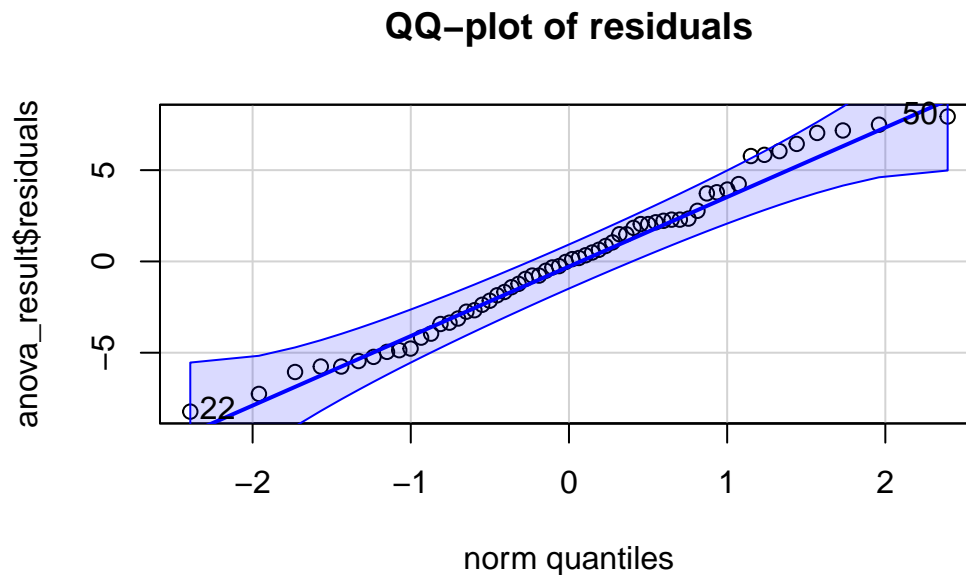
Since all p-values are less than 0.05, we reject all null hypotheses. Therefore, there is sufficient evidence to conclude that the dose of vitamin C, delivery method, and their interaction have significant effects on tooth growth in guinea pigs.

d

```
library(car)
```

Loading required package: carData

```
qqPlot(anova_result$residuals, main = "QQ-plot of residuals")
```



```
[1] 22 50
```

Non-parametric tests

Application of testing the goodness of fit

Testing whether there is a “good fit” between the observed data and the assumed probability model amounts to testing:

Construction of test statistics with an example of 2 categories

Population is 60% female and 40% male. Then, if a sample of 100 students yields 53 females and 47 males, can we conclude that the sample is (random and) representative of the population? That is, how “good” do the data “fit” the assumed probability model of 60% female and 40% male?

Here, let Y_1 denote the number of females selected, $Y_1 \sim B(n, p_1)$ and let Y_2 denote males selected, $Y_2 = (n - Y_1) \sim B(n, p_2) = B(n, 1 - p_1)$.

for samples satisfying the general rule of thumb (the expected number of successes must be at least 5 and the expected number of failures must be at least 5), we can use the normal approximation to the binomial distribution. The test statistic is given by

$$Z = \frac{Y_1 - np_1}{\sqrt{np_1(1 - p_1)}} \sim N(0, 1)$$

which is at least approximately normally distributed.

and

$$Z^2 = Q_1 = \frac{(Y_1 - np_1)^2}{np_1(1 - p_1)} \sim \chi^2(1)$$

which is an approximate chi-square distribution with one degree of freedom.

Now we can multiply Q_1 by $1 = (1 - p_1) + p_1$ to get

$$Q_1 = \frac{(Y_1 - np_1)^2(1 - p_1)}{np_1(1 - p_1)} + \frac{(Y_1 - np_1)^2 p_1}{np_1(1 - p_1)} \sim \chi^2(1)$$

Since $Y_1 = n - Y_2$ and $p_1 = 1 - p_2$, after simplifying, we have

$$Q_1 = \frac{(Y_1 - np_1)^2}{np_1} + \frac{(-(Y_2 - np_2))^2}{np_2} \sim \chi^2(1)$$

which is $Q_1 = \sum_{i=1}^2 \frac{(Y_i - np_i)^2}{np_i} = \sum_{i=1}^2 \frac{(\text{Observed} - \text{Expected})^2}{\text{Expected}} \sim \chi^2(1)$

Hence, it is observed that if the observed counts are very different from the expected counts, then the test statistic will be large. So we reject the null hypothesis if Q_1 is large and how large is large is determined by the critical value of the chi-square distribution with one degree of freedom.

The statistics Q_1 is called the chi-square goodness of fit statistic.

Going back to the example,

- $H_0: p_F = 0.6$
- $H_1: p_F \neq 0.6$

we can calculate the test statistic using a significant level of $\alpha = 0.05$ ($\chi_{0.05,1}^2 = 3.84$) as follows:

$$Q_1 = \frac{(53 - 60)^2}{60} + \frac{(47 - 40)^2}{40} = 2.04$$

Since $Q_1 = 2.04 < 3.84$, we do not reject the null hypothesis. Therefore, we conclude that the sample is (random and) representative of the population.

This can be extended to k categories

Construction of test statistics with an example of k categories

For categories more than 2, i.e.

Categories	1	2	...	$k - 1$	k
Observed	Y_1	Y_2	...	Y_{k-1}	$n - Y_1 - Y_2 - \dots - Y_{k-1}$
Expected	np_1	np_2	...	np_{k-1}	np_k

Karl Pearson showed that the chi-square statistic Q_{k-1} defined as:

$$Q_{k-1} = \sum_{i=1}^k \frac{(Y_i - np_i)^2}{np_i}$$

follows approximately a chi-square random variable with $k - 1$ degrees of freedom. Let's try it out on an example.

- Example:

Categories	Brown	Yellow	Orange	Green	Coffee	Total
Observed y_i	224	119	130	48	59	580
Assumed $H_0(p_i)$	0.4	0.2	0.2	0.1	0.1	1.0
Expected np_i	232	116	116	58	58	580

$$Q_4 = \frac{(224-232)^2}{232} + \frac{(119-116)^2}{116} + \frac{(130-116)^2}{116} + \frac{(48-58)^2}{58} + \frac{(59-58)^2}{58} = 3.784$$

Because there are $k = 5$ categories, we have to compare our chisquare statistic Q_4 to a chi-square distribution with $k - 1 = 5 - 1 = 4$ degrees of freedom:

$$Q_4 = 3.784 < \chi_{4,0.05}^2 = 9.488$$

we fail to reject the null hypothesis.

Application of testing for homogeneity

This is to look at a method for testing whether two or more multinomial distributions are equal.

- Example:

Test the hypothesis that the acceptances of males and females are ditributed equally among the four schools,

(Acceptances)	Bus	Eng	L Arts	Sci	(FIXED) Total
Male	240 (20%)	480 (40%)	120 (10%)	360 (30%)	1200
Female	240 (30%)	80 (10%)	320 (40%)	160 (20%)	800
Total	480 (24%)	560 (28%)	440 (22%)	520 (26%)	2000

Here,

$$H_0 : p_{MB} = p_{FB}, p_{ME} = p_{FE}, p_{ML} = p_{FL}, \text{ and } p_{MS} = p_{FS}$$

$$H_1 : p_{MB} \neq p_{FB} \text{ or } p_{ME} \neq p_{FE} \text{ or } p_{ML} \neq p_{FL}, \text{ or } p_{MS} \neq p_{FS}$$

where:

- p_{Mj} is the proportion of males accepted into school $j = B, E, L, S$.
- p_{Fj} is the proportion of females accepted into school $j = B, E, L, S$.

In conducting such a hypothesis test, we're comparing the proportions of two multinomial distributions.

#(Acc)	Bus ($j = 1$)	Eng ($j = 2$)	L Arts ($j = 3$)	Sci ($j = 4$)	(FIXED) Total
M($i = 1$)	$y_{11} (\hat{p}_{11})$	$y_{12} (\hat{p}_{12})$	$y_{13} (\hat{p}_{13})$	$y_{14} (\hat{p}_{14})$	$n_1 = \sum_{j=1}^k y_{1j}$
F ($i = 2$)	$y_{21} (\hat{p}_{21})$	$y_{22} (\hat{p}_{22})$	$y_{23} (\hat{p}_{23})$	$y_{24} (\hat{p}_{24})$	$n_2 = \sum_{j=1}^k y_{2j}$
Total	$y_{11} + y_{21} (\hat{p}_1)$	$y_{12} + y_{22} (\hat{p}_2)$	$y_{13} + y_{23} (\hat{p}_3)$	$y_{14} + y_{24} (\hat{p}_4)$	$n_1 + n_2$

The chi-square test statistic for testing the equality of two multinomial distributions:

$$Q = \sum_{i=1}^2 \sum_{j=1}^k \frac{(y_{ij} - n_i \hat{p}_j)^2}{n_i \hat{p}_j}$$

follows an approximate chi-square distribution with $k - 1$ degrees of freedom. Reject the null hypothesis of equal proportions if Q is large (since if male and female distributed nearly equally, the expected number of each should be $n_i \hat{p}_j$):

$$Q \geq \chi_{\alpha, k-1}^2$$

(omit the derive of the above Q)

Generally,

$$Q = \sum_{i=1}^h \sum_{j=1}^k \frac{(y_{ij} - n_i \hat{p}_j)^2}{n_i \hat{p}_j} \sim \chi_{(h-1)(k-1)}^2$$

Further example

The head of a surgery department at a university medical center was concerned that surgical residents in training applied unnecessary blood transfusions at a different rate than the more experienced attending physicians. Therefore, he ordered a study of the 49 Attending Physicians and 71 Residents in Training with privileges at the hospital. For each of the 120 surgeons, the number of blood transfusions prescribed unnecessarily in a one-year period was recorded. Based on the number recorded, a surgeon was identified as either prescribing unnecessary blood transfusions Frequently, Occasionally, Rarely, or Never. Here's a summary table (or "contingency table") of the resulting data:

Physician	Frequent	Occasionally	Rarely	Never	Total
Attending	6.942	12.658	22.05	7.35	49
Resident	10.058	18.342	31.95	10.65	71
Total	17	31	54	18	120

Here,

$$H_0 : p_{RF} = p_{AF}, p_{RO} = p_{AO}, p_{RR} = p_{AR}, \text{ and } p_{RN} = p_{AN}$$

$$H_1 : p_{RF} \neq p_{AF} \text{ or } p_{RO} \neq p_{AO} \text{ or } p_{RR} \neq p_{AR}, \text{ or } p_{RN} \neq p_{AN}$$

We should also calculate the expected counts under the null hypothesis. The expected counts are calculated as follows:

Physician	Frequent	Occasionally	Rarely	Never	Total
Attending	6.942	12.658	22.05	7.35	49
Resident	10.058	18.342	31.95	10.65	71
Total	17	31	54	18	120

where, for example, $6.942 = \frac{17}{120} \times 49$ and $10.058 = \frac{17}{120} \times 71$. Now that we have the observed and expected counts, calculating the chisquare statistic is a straightforward exercise:

$$Q = \frac{(2 - 6.942)^2}{6.942} + \dots + \frac{(5 - 10.65)^2}{10.65} = 31.88$$

The chi-square test tells us to reject the null hypothesis, at the 0.05 level, if Q is greater than a chi-square random variable with 3 degrees of freedom, that is, if $Q = 31.88 > 7.815$, we reject the null hypothesis.

Application of testing for independence

This is to look at whether two or more categorical variables are independent.

(previously, the sampling scheme involves: Taking two random (and therefore independent) samples with n_1 and n_2 fixed in advance and observing into which of the k categories the first random samples fall, and observing into which of the k categories the second random samples fall.)

lets consider a different example to illustrate an alternative sampling scheme. Suppose 395 people are randomly selected, and are “cross-classified” into one of eight cells, depending into which age category they fall and whether or not they support legalizing marijuana:

(the sampling scheme involves: Taking one random sample of size n , with n fixed in advance, and then “cross-classifying” each subject into one and only one of the mutually exclusive and exhaustive $A_i \cap B_j$ cells.)

Marijuana Support		Variable B (Age)				
Variable A	OBSERVED	$(18-24)B_1$	$(25-34)B_2$	$(35-49)B_3$	$(50-64)B_4$	Total
(YES) A_1	60	54	46	41		201
(NO) A_2	40	44	53	57		194
Total	100	98	99	98		$n = 395$

Here,

H_0 : Variable A is independent of variable B, that is $P(A_i \cap B_j) = P A_i \times P B_j$ for all i and j

H_1 : Variable A is not independent of variable B.

Generally,

Suppose we have k (column) levels of Variable B indexed by the letter j , and h (row) levels of Variable A indexed by the letter i . Then, we can summarize the data and probability model in tabular format, as follows:

Variable B					
Variable A	$B_1(j = 1)$	$B_2(j = 2)$	$B_3(j = 3)$	$B_4(j = 4)$	Total
$A_1(i = 1)$	$Y_{11}(p_{11})$	$Y_{12}(p_{12})$	$Y_{13}(p_{13})$	$Y_{14}(p_{14})$	$(p_{1.})$
$A_2(i = 2)$	$Y_{21}(p_{21})$	$Y_{22}(p_{22})$	$Y_{23}(p_{23})$	$Y_{24}(p_{24})$	$(p_{2.})$
Total	$(p_{.1})$	$(p_{.2})$	$(p_{.3})$	$(p_{.4})$	n

where $p_{ij} = Y_{ij}/n$, $p_{i.} = \sum_{j=1}^k p_{ij}$, and $p_{.j} = \sum_{i=1}^h p_{ij}$

$$Q = \sum_{j=1}^k \sum_{i=1}^h \frac{(y_{ij} - \frac{y_{i.} \cdot y_{.j}}{n})^2}{\frac{y_{i.} \cdot y_{.j}}{n}} \sim \chi_{(h-1)(k-1)}^2$$

Are chi-square statistic for homogeneity and the chi-square statistic for independence equivalent?

Although their chi-square statistics are equivalent, the two tests are not equivalent since their sampling experiment designs are different.

Here’s the table of expected counts:

Bicycle Riding Interest		Variable B (Age)				
Variable A	EXPECTED	18-24	25-34	35-49	50-64	Total
	YES	50.886	49.868	50.377	49.868	201
	NO	49.114	48.132	48.623	48.132	194
	Total	100	98	99	98	395

$$Q = \frac{(60 - 50.886)^2}{50.886} + \dots + \frac{(57 - 48.132)^2}{48.132} = 8.006$$

The chi-square test tells us to reject the null hypothesis, at the 0.05 level, since Q is greater than a chi-square random variable with 3 degrees of freedom, that is, $Q = 8.006 > 7.815$.

Summary

Parametric tests make assumptions that aspects of the data follow some sort of theoretical probability distribution. Non-parametric tests or distribution free methods do not, and are used when the distributional assumptions for a parametric test are not met. While this is an advantage, it often comes at a cost of power (in the sense they are less likely to be able to detect a difference when a true difference exists).

Most non-parametric tests are just hypothesis tests; there is no estimation of a confidence interval.

Most non-parametric methods are based on ranking the values of a variable in ascending order and then calculating a test statistic based on the sums of these ranks.

Non-parametric tests include:

- Two-sample independent t-test - Wilcoxon rank-sum test or Mann-Whitney U test
- Paired t-test - Wilcoxon signed-rank test
- One-way ANOVA - Kruskal-Wallis Test
- Normality tests - Shapiro-Wilk test and Kolmogorov-Smirnov test

Linear Regression

Simple linear regression

Brief introduction

Some people think simple methods is bad and like complicated methods , but actually simple is very good-SLR works very well in lots of situations.

SLR is used to answer:

is there a relationship between..

How strong the relationship is

which variable contribute to this relationship

How accurate could we predict the response variable

Is the relationship linear?

Is there a synergy among independent variables?

This is a model with two random variables, X and Y, where we are trying to predict Y from X. Here are the model's assumptions:

- The distribution of X is arbitrary, possibly is even non-random;
- If $X=x$, then $Y = \beta_0 + \beta_1 x + \epsilon$ for some constants β_0, β_1 and some random noise variable ϵ
- ϵ has mean 0, a constant variance σ^2 , and is uncorrelated with X and uncorrelated across observations $\text{Cov}(\epsilon_i, \epsilon_j) = 0$ for $i \neq j$

Using Least Squares, we can estimate $\hat{\beta}_0$ and $\hat{\beta}_1$, which are unbiased estimates of β_0 and β_1 .

- Gaussian-Noise Simple Linear Regression Model

Now we further assume that the distribution of ϵ is normal, i.e. $\epsilon \sim N(0, \sigma^2)$, independent of X.

They tell us, exactly, the probability distribution for Y given X, and so will let us get exact distributions for predictions and for other inferential statistics.

Maximum Likelihood Estimation (MLE)

Introduction to MLE

Likelihood is a fundamental concept in statistics that measures how well a particular set of parameters (e.g., the mean of a distribution) explains observed data. Think of it as a “score” that tells you which parameter values make your data most plausible.

Compared to probability, which answers: “What’s the chance of seeing this data if we assume specific parameters?” , likelihood answers: “Given this data, how plausible are these parameters?”

If the parameters are b_0, b_1, s^2 (reserving the Greek letters for their true values), then $Y | X = x \sim N(b_0 + b_1 x, s^2)$, and Y_i and Y_j are independent given X_i and X_j , so the overall likelihood is

$$\prod_{i=1}^n \frac{1}{\sqrt{2\pi s^2}} e^{-\frac{(y_i - (b_0 + b_1 x_i))^2}{2s^2}}$$

As usual, we work with the log-likelihood, which gives us the same information but replaces products with sums:

$$L(b_0, b_1, s^2) = -\frac{n}{2} \ln(2\pi s^2) - \frac{1}{2s^2} \sum_{i=1}^n (y_i - (b_0 + b_1 x_i))^2$$

maximize it:

$$\frac{\partial L}{\partial b_0} = -\frac{1}{2s^2} \sum_{i=1}^n 2(y_i - (b_0 + b_1 x_i))(-1)$$

$$\frac{\partial L}{\partial b_1} = -\frac{1}{2s^2} \sum_{i=1}^n 2(y_i - (b_0 + b_1 x_i))(-x_i)$$

Same result of MLE as least squares in linear regression

Notice that when we set these derivatives to zero, all the multiplicative constants - in particular, the prefactor of $1/2s^2$ - go away. We are left with

$$\begin{aligned} \sum_{i=1}^n (y_i - (\hat{\beta}_0 + \hat{\beta}_1 x_i)) &= 0 \\ \sum_{i=1}^n (y_i - (\hat{\beta}_0 + \hat{\beta}_1 x_i)) x_i &= 0 \end{aligned}$$

These are, up to a factor of $1/n$, exactly the equations we got from the method of least squares. That means that the least squares solution is the maximum likelihood estimate under the Gaussian noise model.

Maximum likelihood estimates of the regression curve coincide with least-squares estimates when the noise around the curve is additive, Gaussian, of constant variance, and both independent of X and of other noise terms. If any of those assumptions fail, maximum likelihood and least squares estimates can diverge.

Hypothesis testing for estimates with unknown σ^2

Residual sum of squares (RSS) and t-statistics construction in linear regression

It can be shown that (the proof is beyond the scope of this course)

$$\frac{RSS}{\sigma^2} = \frac{(n-2)\hat{\sigma}^2}{\sigma^2} \sim \chi_{n-2}^2$$

This allows us to construct a t -value

$$t = \frac{\hat{\beta} - \beta}{s_{\hat{\beta}}} \sim t_{n-2}$$

Under the normality assumption of the error terms, the estimator of the slope coefficient will itself be normally distributed with mean β_i and variance $\text{Var}[\beta_i]$. For $\hat{\beta}_1$, its mean is β_1 and its variance is $\sigma^2 / \sum (x_i - \bar{x})^2$. When σ^2 is known, we know

$$\frac{\hat{\beta}_1 - \beta_1}{\frac{\sigma}{\sqrt{\sum_{i=1}^n (x_i - \bar{x})^2}}}$$

follows standard normal distribution. However, in practice, σ^2 is often unknown. We then divide this standard normal distributed term by

$$\sqrt{\frac{(n-2)\hat{\sigma}^2}{(n-2)\sigma^2}} = \frac{\hat{\sigma}}{\sigma}$$

Therefore, when we write

$$s_{\hat{\beta}_1} = \frac{\hat{\sigma}}{\sqrt{\sum_{i=1}^n (x_i - \bar{x})^2}}$$

we construct a t -statistic for $\hat{\beta}_1$ with degrees of freedom $n-2$. This then allows us to construct a $100(1-\alpha)\%$ confidence interval for β_1 :

$$\hat{\beta}_1 \pm t_{n-2, \alpha/2} \times s_{\hat{\beta}_1}$$

We can also do similar calculation to get the t -statistic and confidence interval for β_0 .

Hyphothesis Testing

$$t = \hat{\beta} - \beta / s_{\hat{\beta}_1} \sim t_{n-2}$$

$$H_0 : \beta_i = 0$$

$$H_1 : \beta_i \neq 0$$

Codes of linear regression with CI

```
lmodel <- lm(Petal.Length ~ Petal.Width, data=iris)
summary(lmodel)
```

Call:

```
lm(formula = Petal.Length ~ Petal.Width, data = iris)
```

Residuals:

Min	1Q	Median	3Q	Max
-1.33542	-0.30347	-0.02955	0.25776	1.39453

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	1.08356	0.07297	14.85	<2e-16 ***
Petal.Width	2.22994	0.05140	43.39	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.4782 on 148 degrees of freedom

Multiple R-squared: 0.9271, Adjusted R-squared: 0.9266

F-statistic: 1882 on 1 and 148 DF, p-value: < 2.2e-16

```
confint(lmodel)
```


	2.5 %	97.5 %
(Intercept)	0.9393664	1.227750
Petal.Width	2.1283752	2.331506

```

conf_interval <- predict(lmodel, data=iris, interval='confidence', level=0.95)
plot(iris$Petal.Width, iris$Petal.Length,
     xlab='Petal.Width', ylab='Petal.Length',
     main='Simple Linear Regression')
abline(lmodel, col='lightblue')
matlines(iris$Petal.Width, conf_interval[,2:3], col='blue', lty=2)

```

Simple Linear Regression




```

# Using ggplot2
library(ggplot2)
ggplot(iris, aes(x=Petal.Width, y=Petal.Length)) +
  geom_point() +
  geom_smooth(method=stats::lm, se=T, level=0.95)

```

`geom_smooth()` using formula = 'y ~ x'



Linear Regression and ANOVA

```
fev_dat <- read.table('fev_dat.txt', header=T)
fev_dat_subset <- fev_dat[fev_dat$age >= 6 & fev_dat$age <= 10,]
ggplot(fev_dat_subset, aes(x=age, y=FEV)) +
  geom_point() +
  geom_smooth(method=stats::lm, se=T, level=0.95)
summary(aov(FEV ~ age, data=fev_dat_subset))
summary(lm(FEV ~ age, data=fev_dat_subset))
anova(lm(FEV ~ age, data=fev_dat_subset))
```

R^2 —the fraction of variability explained by the regression

$$R^2 = 1 - \frac{SSR}{SSTO}$$

Multiple linear regression (MLR)

$y = X\beta + \epsilon$, where y is a $n \times 1$ row vector, X is a $n \times (k + 1)$ matrix, and β is a $(k + 1) \times 1$ column vector for all n observations.

A potential problem in practice –multicollinearity

When multicollinearity exists, any of the following pitfalls can be exacerbated:

- The estimated regression coefficient of any one variable depends on which other predictors are included in the model
- The precision of the estimated regression coefficients decreases as more predictors are added to the model
- The marginal contribution of any one predictor variable in reducing the error sum of squares depends on which other predictors are already in the model
- Hypothesis tests for $\beta_j = 0$ may yield different conclusions depending on which predictors are in the model

Perfect multicollinearity

Perfect multicollinearity refers to a situation where the predictive variables have an exact linear relationship. When there is perfect collinearity, the design matrix X has less than full rank, and therefore the moment matrix $X'X$ cannot be inverted. In this situation, the parameter estimates of the regression are not well-defined, as the system of equations has infinitely many solutions.

Imperfect multicollinearity

Imperfect multicollinearity refers to a situation where the predictive variables have a nearly exact linear relationship.

$R^2 = r^2$ where r is the Pearson correlation coefficient.

Adjusted R-squared

Adjusted R-squared is a modified version of R-squared that adjusts for the number of predictors in a model. It provides a more accurate measure of the model's explanatory power, penalizing for the addition of irrelevant predictors. This helps in comparing models with different numbers of predictors.

Logistic Regression


This is a regression of categorical outcome.

In regression analysis with a categorical outcome, such as predicting a binary variable (yes or no), simple linear regression is not ideal. This is because:

- The predicted values may fall outside the range of 0 to 1, which is not meaningful for probabilities.
- Small changes in the predictors can lead to relatively small fluctuations in the predicted probabilities near the 0.5 mark (natural threshold), which is actually where decision-making is most critical.

So we need S-curve to satisfy above things.

```
curve(1 / (1 + exp(-x)), from = -6, to = 6, xlab = "x", ylab = "f(x)",  
main = "Logistic Function:  $f(x) = 1 / (1 + \exp(-x))$ ", col = "blue", lwd = 2)
```



An example of a not well-predicted logistic model (since stock is not easy to predict)

```
library(ISLR2)  
attach(Smarket)  
glm.fits <- glm(  
  Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 + Volume,  
  data = Smarket, family = binomial  
)  
summary(glm.fits)
```

Call:

```
glm(formula = Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 +  
     Volume, family = binomial, data = Smarket)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-0.126000	0.240736	-0.523	0.601
Lag1	-0.073074	0.050167	-1.457	0.145
Lag2	-0.042301	0.050086	-0.845	0.398
Lag3	0.011085	0.049939	0.222	0.824
Lag4	0.009359	0.049974	0.187	0.851
Lag5	0.010313	0.049511	0.208	0.835
Volume	0.135441	0.158360	0.855	0.392

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1731.2 on 1249 degrees of freedom
Residual deviance: 1727.6 on 1243 degrees of freedom
AIC: 1741.6

Number of Fisher Scoring iterations: 3

```
glm.probs <- predict(glm.fits, type = "response")  
glm.probs[1:10]
```

1	2	3	4	5	6	7	8
0.5070841	0.4814679	0.4811388	0.5152224	0.5107812	0.5069565	0.4926509	0.5092292
9	10						
0.5176135	0.4888378						

```
glm.pred <- rep("Down", 1250)  
glm.pred[glm.probs > .5] = "Up"  
table(glm.pred, Direction)
```

	Direction	
glm.pred	Down	Up
Down	145	141
Up	457	507

```
(507 + 145) / 1250
```

```
[1] 0.5216
```

```
mean(glm.pred == Direction)
```

```
[1] 0.5216
```

Odds

Odds are another way of quantifying the probability of an event, commonly used in gambling (and logistic regression).

For some event E ,

$$\text{odds}(E) = \frac{P(E)}{P(E^c)} = \frac{P(E)}{1 - P(E)}$$

The odds ratio (OR) is the ratio of the odds of an event occurring in one group to the odds of it occurring in another group. If the event in each of the groups are p_1 (first group) and p_2 (second group), then the odds ratio is:

$$\text{OR} = \frac{p_1 / (1 - p_1)}{p_2 / (1 - p_2)}$$

Generalized Linear Models

Under the hood, we're still using a linear model ($\beta_0 + \beta_1 x$), but now it's embedded in a function that ensures valid probabilities. This is the essence of logistic regression - a generalized linear model (GLM) designed for binary outcomes.

Given predictor X and an outcome Y , a GLM is defined by three components: - A random component, that specifies a distribution for $Y | X$ - A systematic component, that relates a parameter η to the predictor X

$$\eta = \beta_0 + \beta_1 X_1 + \dots + \beta_k X_k$$

- A link function, that connects the random and systematic component

Random Component

The random component specifies a distribution for the outcome variable (conditional on X). In the case of linear regression, we assume that $Y | X \sim \mathcal{N}(\mu, \sigma^2)$, for some mean μ and variance σ^2 . In the case of logistic regression, we assume that $Y | X \sim \text{Bern}(p)$ for some probability p .

In a generalized model, we are allowed to assume that $Y | X$ has a probability density function or probability mass function of the form

$$f(y; \theta, \phi) = \exp \left(\frac{y\theta - b(\theta)}{a(\phi)} + c(y, \phi) \right)$$

Here θ, ϕ are parameters, and a, b, c are functions. Any density of the above form is called an exponential family density. The parameter θ is called the natural parameter, and the parameter ϕ the dispersion parameter.

Exponential Family

Exponential families include many of the most common distributions. For example: - Exponential

$$f(y; \lambda) = \lambda e^{-\lambda y} = \exp(-y\lambda + \ln \lambda)$$

where $\theta = -\lambda, \phi = 1, b(\theta) = \ln \lambda, a(\phi) = 1$, and $c(y, \phi) = 0$ - Poisson

$$f(y; \lambda) = \frac{e^{-\lambda} \lambda^y}{y!} = \exp(y \ln \lambda - \lambda - \ln(y!))$$

where $\theta = \ln \lambda, \phi = 1, b(\theta) = e^\theta = \lambda, a(\phi) = 1$, and $c(y, \phi) = -\lambda - \ln(y!)$

Systematic Component and Link Component

The systematic component relates a parameter η to the predictors X . In a GLM, this is always done via

$$\eta = X\beta = \beta_0 + \beta_1 X_1 + \dots + \beta_k X_k$$

We will denote the expectation of the distribution in random component as μ , i.e., $\mathbb{E}[Y | X] = \mu$. It will be our goal to estimate μ . Finally, the link component connects the random and systematic components, via a link function g . In particular, this link function provides a connection between μ and η , as in

$$g(\mu) = \eta \quad \text{or} \quad \mu = g^{-1}(\eta)$$

Example

Gaussian-noise Linear Regression

- Random Component: $Y | X \sim \mathcal{N}(\mu, \sigma^2)$ and $\mathbb{E}[Y | X] = \mu$
- Systematic Component: $\eta = X\beta$
- Link Component: $g(\mu) = \mu$, so that $\mu = \eta = X\beta$

Bernoulli

Suppose that $Y \in \{0, 1\}$, and we model the distribution of $Y | X$ as Bernoulli with success probability p . Then the probability mass function (not a density, since Y is discrete) is

$$f(y) = p^y(1 - p)^{1-y}$$

We can rewrite to fit the exponential family form as

$$\begin{aligned} f(y) &= \exp(y \log p + (1 - y) \log(1 - p)) \\ &= \exp(y \log(p/(1 - p)) + \log(1 - p)) \end{aligned}$$

$$f(y; \theta, \phi) = \exp \left(\frac{y\theta - b(\theta)}{a(\phi)} + c(y, \phi) \right)$$

Here we would identify $\theta = \log(p/(1 - p))$ as the natural parameter. Note that the mean here is $\mu = p$, and using the inverse of the above relationship, we can directly write the mean p as a function of θ , as in $p = e^\theta / (1 + e^\theta)$. Hence $b(\theta) = \log(1 - p) = -\log(1 + e^\theta)$. There is no dispersion parameter, so we can set $a(\phi) = 1$. Also, $c(y, \phi) = 0$.

Link Function

$$g(\mu) = \Phi^{-1}(\mu)$$

where Φ is the standard normal CDF.

Logistic

The three GLM criteria give us:

- $y_i \sim \text{Bern}(p_i)$
- $\eta = \beta_0 + \beta_1 X_1 + \dots + \beta_k X_k$
- $\text{logit}(p) = \log \frac{p}{1-p} = \eta$

From which we know,

$$p_i = \frac{\exp(\beta_0 + \beta_1 x_{i1} + \dots + \beta_k X_{ik})}{1 + \exp(\beta_0 + \beta_1 x_{i1} + \dots + \beta_k X_{ik})}$$

Survival Analysis

Survival analysis is used to analyze data in which the time until the event is of interest. The response is often referred to as a failure time, survival time, or event time.

Some definitions

Hazard ratios; ratios of hazard functions between different groups (e.g., exposed vs. unexposed) while adjusting for confounders.

censoring - which occurs when the survival time is only partially known

- Fixed type I censoring occurs when a study is designed to end after C years of follow-up. In this case, everyone who does not have an event observed during the course of the study is censored at C years.
- In random type I censoring, the study is designed to end after C years, but censored subjects do not all have the same censoring time. This is the main type of right-censoring we will be concerned with.
- In type II censoring, a study ends when there is a pre-specified number of events.

Kaplan-Meier estimate


```
library(ISLR2)
names(BrainCancer)
```

```
[1] "sex"      "diagnosis" "loc"      "ki"      "gtv"      "stereo"
[7] "status"   "time"
```

```
attach(BrainCancer)
table(status)
```

```
status
 0  1
53 35
```

```
library(survival)
fit.surv <- survfit(Surv(time, status) ~ 1)
plot(fit.surv, xlab = "Months",
     ylab = "Estimated Probability of Survival")
```



Cox-proportional hazards model

$$S(t) = P(T > t) = 1 - F(t)$$

$$h(t) = \lim_{\Delta t \rightarrow 0} \frac{P(t < T \leq t + \Delta t | T > t)}{\Delta t}$$


```
fit.all <- coxph(
Surv(time, status) ~ sex + diagnosis + loc + ki + gtv +
  stereo)
fit.all
```

Call:

```
coxph(formula = Surv(time, status) ~ sex + diagnosis + loc +
      ki + gtv + stereo)
```


	coef	exp(coef)	se(coef)	z	p
sexMale	0.18375	1.20171	0.36036	0.510	0.61012
diagnosisLG glioma	0.91502	2.49683	0.63816	1.434	0.15161
diagnosisHG glioma	2.15457	8.62414	0.45052	4.782	1.73e-06
diagnosisOther	0.88570	2.42467	0.65787	1.346	0.17821
locSupratentorial	0.44119	1.55456	0.70367	0.627	0.53066
ki	-0.05496	0.94653	0.01831	-3.001	0.00269
gtv	0.03429	1.03489	0.02233	1.536	0.12466
stereoSRT	0.17778	1.19456	0.60158	0.296	0.76760

Likelihood ratio test=41.37 on 8 df, p=1.776e-06

n= 87, number of events= 35

(1 observation deleted due to missingness)

```
modaldata <- data.frame(
  diagnosis = levels(diagnosis),
  sex = rep("Female", 4),
  loc = rep("Supratentorial", 4),
  ki = rep(mean(ki), 4),
  gtv = rep(mean(gtv), 4),
  stereo = rep("SRT", 4)
)
survplots <- survfit(fit.all, newdata = modaldata)
plot(survplots, xlab = "Months",
      ylab = "Survival Probability", col = 2:5)
legend("bottomleft", levels(diagnosis), col = 2:5, lty = 1)
```



Final review of R codes

Calculation

```
log(exp(1)) # base `e` is the default (log(e) is not defined)
```

```
[1] 1
```

Vectors

```
x <- c(1,2,3,4)
class(x)
```

```
[1] "numeric"
```

```
x %*%x # scalar ("inner") product (but default in R as an 1*1 matrix)
```

```
      [,1]
[1,]    30
```

```
rep(c('a','b'),3)
```

```
[1] "a" "b" "a" "b" "a" "b"
```

```
rep(c(2, 4, 8), each = 3)
```

```
[1] 2 2 2 4 4 4 8 8 8
```

```
y <- seq(from = 1, to = 4, by =1)  
class(x)
```

```
[1] "numeric"
```

```
str(x)
```

```
num [1:4] 1 2 3 4
```

```
x <- c(-5:5)  
str(x)
```

```
int [1:11] -5 -4 -3 -2 -1 0 1 2 3 4 ...
```

```
1:4 # c(1,2,3,4) and 1:4 are the same in.R
```

```
[1] 1 2 3 4
```

```
seq(1,5, length.out=11)
```

```
[1] 1.0 1.4 1.8 2.2 2.6 3.0 3.4 3.8 4.2 4.6 5.0
```

```
# vac<-c((1,2,3),(3,4,5)) is wrong, but the below is true  
vec1 <- c(1,2,3)  
vec2 <- c(4,5,6)  
vec3 <- c(vec1, vec2)  
vec3[1] == vec1[1]
```

```
[1] TRUE
```

```
vec3[3:5];vec3[c(2,3)]
```

```
[1] 3 4 5
```

```
[1] 2 3
```

```
vec3[-1] # everything but the first element
```

```
[1] 2 3 4 5 6
```

```
vec3[-2*c(1,2)]
```

```
[1] 1 3 5 6
```

```
x <- -5:5
```

```
abs(-5:5)
```

```
[1] 5 4 3 2 1 0 1 2 3 4 5
```

```
x <- c(1, 2, 3, 4, 5,6)
y <- c(10,11)
result <- x + y
print(result)
```

```
[1] 11 13 13 15 15 17
```

```
# x * y
c(1,2)*c(2,3)
```

```
[1] 2 6
```

```
c(1,2)%*%c(2,3)
```

```
      [,1]
[1,]      8
```

```
#####Application
# Calculate the sample(var(x)) and population variance
x <- c(1, 2, 3, 4, 5)
n <- length(x)
# Calculate the sample variance using R's var() function
sample_variance <-var(x)
sample_variance
```

```
[1] 2.5
```

```
population_variance <- sample_variance * (n - 1) / n
population_variance
```

```
[1] 2
```

Rmd knowledge

{r, echo = FALSE} –Hidden Code
 {r, eval = FALSE} –Do not run this code
 {r, message = FALSE} –Do not show the message
 {r, warning = FALSE} –Do not show the warning
 {r,results='hide'} –Do not show the results

Probability in R

```
dnorm(0) # density at 0
```

```
[1] 0.3989423
```

```
pnorm(-1) # cumulative probability at -1
```

```
[1] 0.1586553
```

```
pnorm(-1,lower.tail = F) # cumulative probability at -1, upper tail
```

```
[1] 0.8413447
```

```
pnorm(0)
```

```
[1] 0.5
```

```
qnorm(0) # quantile at 0 (with the cumulative probability of 0)
```

```
[1] -Inf
```

```
pnorm(1.645)
```

```
[1] 0.9500151
```

```
qnorm(0.95) # norm quantile at 0.95 (with the cumulative probability of 0.95)
```

```
[1] 1.644854
```

```
pnorm(1.96)
```

```
[1] 0.9750021
```

```
qnorm(0.975)
```

```
[1] 1.959964
```

```
library(mosaic)
```

```
Registered S3 method overwritten by 'mosaic':  
  method from  
  fortify.SpatialPolygonsDataFrame ggplot2
```

The 'mosaic' package masks several functions from core packages in order to add additional features. The original behavior of these functions should not be affected by this.

Attaching package: 'mosaic'

The following objects are masked from 'package:dplyr':

count, do, tally

The following object is masked from 'package:Matrix':

mean

The following objects are masked from 'package:car':

deltaMethod, logit

The following object is masked from 'package:ggplot2':

stat

The following objects are masked from 'package:stats':

binom.test, cor, cor.test, cov, fivenum, IQR, median, prop.test,
quantile, sd, t.test, var

The following objects are masked from 'package:base':

max, mean, min, prod, range, sample, sum

```
ppois(3, lambda = 2)
```

```
[1] 0.8571235
```

```
sum(dpois(0:3, lambda = 2))
```

```
[1] 0.8571235
```

```
ppois(3, lambda = 2) == sum(dpois(0:3, lambda = 2)) # this inequivalence is because of the f
```

```
[1] FALSE
```

R basic

```
paste("Good", "afternoon", "ladies", "and", "gentlemen")
```

```
[1] "Good afternoon ladies and gentlemen"
```

```
paste0("Good", "afternoon", "ladies", "and", "gentlemen")
```

```
[1] "Goodafternoonladiesandgentlemen"
```

```
x <- -10:10  
which(x>0)
```

```
[1] 12 13 14 15 16 17 18 19 20 21
```

```
x[which(x>0)]
```

```
[1] 1 2 3 4 5 6 7 8 9 10
```

```
vowels <- c('a','e','i','o','u')  
which(is.element(letters, vowels))
```

```
[1] 1 5 9 15 21
```

Data class

Numeric

Integer


```
y <- 42L  
class(y)
```

```
[1] "integer"
```

Character

```
x <- "123"  
class(x)
```

```
[1] "character"
```

```
# [1] "character"  
  
x <- as.numeric(x)  
class(x)
```

```
[1] "numeric"
```

```
# [1] "numeric"
```

Logical

```
# Note that logical elements are NOT in quotes.  
z = c("TRUE", "FALSE", "TRUE", "FALSE")  
class(z)
```

```
[1] "character"
```

```
as.logical(z)
```

```
[1] TRUE FALSE TRUE FALSE
```

```
# TRUE = 1 and FALSE = 0. sum() and mean() work on logical vectors
```

```
# remember:
```

```
TRUE & TRUE
```

```
[1] TRUE
```

```
TRUE & FALSE
```

```
[1] FALSE
```

```
TRUE | FALSE
```

```
[1] TRUE
```

Factor

```
y <- c('B','B','A','A','C')
z <- factor(y)
str(z)
```

```
Factor w/ 3 levels "A","B","C": 2 2 1 1 3
```

```
as.numeric(z)
```

```
[1] 2 2 1 1 3
```

```
levels(z)
```

```
[1] "A" "B" "C"
```

```
z <- factor(z,                # vector of data levels to convert
            levels=c('B','A','C'), # Order of the levels
            labels=c("B Group", "A Group", "C Group")) # Pretty labels to use
z
```

```
[1] B Group B Group A Group A Group C Group
Levels: B Group A Group C Group
```

```
#####Application
### eg of use in the plot's x-axis name label

iris$Species <- factor(iris$Species,
                      levels = c('versicolor','setosa','virginica'),
                      labels = c('Versicolor','Setosa','Virginica'))
#boxplot(Sepal.Length ~ Species, data=iris)

### another eg
#age_category <- ifelse(ages >= 18, "Adult", "Minor")
#age_factor <- factor(age_category, levels = c("Minor", "Adult"))
#age_factor

### transform a continuous numerical vector into a factor

x <- 1:10
cut(x, breaks = c(0, 2.5, 5.0, 7.5, 10))
```

```
[1] (0,2.5] (0,2.5] (2.5,5] (2.5,5] (2.5,5] (5,7.5] (5,7.5] (7.5,10]
[9] (7.5,10] (7.5,10]
Levels: (0,2.5] (2.5,5] (5,7.5] (7.5,10]
```

```
x<-cut(x, breaks=3, labels=c('Low','Medium','High'))
str(x)
```

```
Factor w/ 3 levels "Low","Medium",...: 1 1 1 1 2 2 2 3 3 3
```

Date and time

The following symbols can be used with the `format()` function to print dates.

- %d day as a number (0-31) 01-31
- %a abbreviated weekday Mon
- %A unabbreviated weekday Monday
- %m month (00-12) 00-12
- %b abbreviated month Jan
- %B unabbreviated month January
- %y 2-digit year 07
- %Y 4-digit year 2007

```
mydates <- as.Date(c("2023-04-07", "2023-01-01"))
mydates
```

```
[1] "2023-04-07" "2023-01-01"
```

```
days <- mydates[1] - mydates[2]; days
```

Time difference of 96 days

```
today <- Sys.Date()
format(today, format="%B %d %Y")
```

```
[1] "June 08 2025"
```

Data frame

```
Data_Frame <- data.frame(Tr =c("1","2","3"),
                          Pu =c(11,21,32),
                          Dur=c(22,222,1))
Data_Frame
```

```
  Tr Pu Dur
1  1 11  22
2  2 21 222
3  3 32   1
```

```
summary(Data_Frame)
```

	Tr	Pu	Dur
Length:	3	Min. :11.00	Min. : 1.00
Class :	character	1st Qu.:16.00	1st Qu.: 11.50
Mode :	character	Median :21.00	Median : 22.00
		Mean :21.33	Mean : 81.67
		3rd Qu.:26.50	3rd Qu.:122.00
		Max. :32.00	Max. :222.00

```
### table and data frame
table(mpg$class)
```

```

      2seater   compact   midsize   minivan   pickup subcompact      suv
         5         47         41         11         33         35         62

```

```
df <- as.data.frame(table(mpg$class))
df
```

```

      Var1 Freq
1   2seater    5
2   compact   47
3   midsize   41
4   minivan   11
5   pickup    33
6 subcompact   35
7      suv     62

```

List

```

# List ---Can hold vectors, strings, matrices, models, list of other list, lists upon lists!

mylist <- list(letters=c("a","b","c"),
               numbers=1:3,matrix(1:25,ncol=5))
head(mylist)

```

```

$letters
[1] "a" "b" "c"

```

```

$numbers
[1] 1 2 3

```

```

[[3]]
      [,1] [,2] [,3] [,4] [,5]
[1,]    1    6   11   16   21
[2,]    2    7   12   17   22
[3,]    3    8   13   18   23
[4,]    4    9   14   19   24
[5,]    5   10   15   20   25

```

```
# Can reference data using $ (if the elements are named), or using [], or [[]]
```

```
mylist[1] # list
```

```
$letters  
[1] "a" "b" "c"
```

```
mylist["letters"] # list
```

```
$letters  
[1] "a" "b" "c"
```

```
mylist[[1]] # vector
```

```
[1] "a" "b" "c"
```

```
mylist$letters == mylist[["letters"]]
```

```
[1] TRUE TRUE TRUE
```

```
mylist[[3]][1:2,1:2]
```

```
      [,1] [,2]  
[1,]    1    6  
[2,]    2    7
```

```
class(mylist[[3]][1:2,1:2])
```

```
[1] "matrix" "array"
```

```
x = c(0, 2, 2, 3, 4); 2 %in% x
```

```
[1] TRUE
```

```
# eg of using list
x <- c(5.1, 4.9, 5.6, 4.2, 4.8, 4.5, 5.3, 5.2) # some toy data
results <- t.test(x, alternative='less', mu=5) # do a t-test
str(results)
```

List of 10

```
$ statistic : Named num -0.314
..- attr(*, "names")= chr "t"
$ parameter : Named num 7
..- attr(*, "names")= chr "df"
$ p.value : num 0.381
$ conf.int : num [1:2] -Inf 5.25
..- attr(*, "conf.level")= num 0.95
$ estimate : Named num 4.95
..- attr(*, "names")= chr "mean of x"
$ null.value : Named num 5
..- attr(*, "names")= chr "mean"
$ stderr : num 0.159
$ alternative: chr "less"
$ method : chr "One Sample t-test"
$ data.name : chr "x"
- attr(*, "class")= chr "htest"
```

```
results$p.value
```

```
[1] 0.3813385
```

Matrix

```
# Matrices

n=1:9
mat = matrix(n,nrow=3)
mat
```

```
      [,1] [,2] [,3]
[1,]    1    4    7
[2,]    2    5    8
[3,]    3    6    9
```

```

y <- diag(n)

# use %*% as the product of matrices

## Eigenvalue and Eigenvector

A <- matrix(c(13, -4, 2, -4, 11, -2, 2, -2, 8), 3, 3, byrow=TRUE)
ev <- eigen(A)

(values <- ev$values)

```

```
[1] 17  8  7
```

```
(vectors <- ev$vectors)
```

```

           [,1]      [,2]      [,3]
[1,] 0.7453560 0.6666667 0.0000000
[2,] -0.5962848 0.6666667 0.4472136
[3,] 0.2981424 -0.3333333 0.8944272

```

```
## Data selection --row then column
```

```
mat[1, 1]
```

```
[1] 1
```

```
mat[1,]
```

```
[1] 1 4 7
```

```
mat[,1]
```

```
[1] 1 2 3
```

```
class(mat[1, ]) # Note that the class of the returned object is no longer a matrix
```

```
[1] "integer"
```


Example 1

```
db_data <- list(  
  drugs = list(  
    general_information = data.frame(  
      drugbank_id = c("DB001", "DB002", "DB003", "DB004", "DB005"),  
      name = c("Aspirin", "Ibuprofen", "Paracetamol", "Insulin", "Morphine"),  
      type = c("small molecule", "small molecule", "small molecule", "biotech", "small molecule"),  
      created = as.Date(c("2020-01-01", "2020-02-01", "2020-03-01", "2020-04-01", "2020-05-01")),  
      stringsAsFactors = FALSE  
    ),  
    drug_classification = data.frame(  
      drugbank_id = c("DB001", "DB002", "DB003", "DB004", "DB005"),  
      classification = c("Analgesic", "Anti-inflammatory", "Analgesic", "Hormone", "Analgesic"),  
      stringsAsFactors = FALSE  
    ),  
    experimental_properties = data.frame(  
      drugbank_id = c("DB001", "DB002", "DB003", "DB004", "DB005", "DB001", "DB002", "DB003",  
        kind = c("logP", "logP", "logP", "logP", "logP", "Molecular Weight", "Molecular Weight",  
        value = c("1.2", "1.5", "0.8", "2.1", "1.8", "180.1", "206.3", "151.2", "5800.0", "281.4",  
        stringsAsFactors = FALSE  
    )  
  )  
)  
db_data
```

```
$drugs  
$drugs$general_information  
  drugbank_id      name      type      created  
1      DB001    Aspirin small molecule 2020-01-01  
2      DB002  Ibuprofen small molecule 2020-02-01  
3      DB003 Paracetamol small molecule 2020-03-01  
4      DB004    Insulin      biotech 2020-04-01  
5      DB005    Morphine small molecule 2020-05-01
```

```
$drugs$drug_classification  
  drugbank_id      classification  
1      DB001          Analgesic  
2      DB002 Anti-inflammatory  
3      DB003          Analgesic  
4      DB004          Hormone
```

```
5          DB005          Analgesic
```

```
$drugs$experimental_properties
  drugbank_id      kind  value
1      DB001      logP    1.2
2      DB002      logP    1.5
3      DB003      logP    0.8
4      DB004      logP    2.1
5      DB005      logP    1.8
6      DB001 Molecular Weight 180.1
7      DB002 Molecular Weight 206.3
8      DB003 Molecular Weight 151.2
9      DB004 Molecular Weight 5800.0
10     DB005 Molecular Weight 281.5
```

```
general_information <- db_data$drugs$general_information
print(general_information)
```

```
  drugbank_id      name      type  created
1      DB001  Aspirin small molecule 2020-01-01
2      DB002  Ibuprofen small molecule 2020-02-01
3      DB003 Paracetamol small molecule 2020-03-01
4      DB004   Insulin      biotech 2020-04-01
5      DB005  Morphine small molecule 2020-05-01
```

```
# 20. Number of drugs in the general_information dataframe
general_information <- db_data$drugs$general_information
nrow(general_information)
```

```
[1] 5
```

```
# 21. Filter drugs of type "biotech"
general_information[general_information$type == 'biotech',]
```

```
  drugbank_id      name      type  created
4      DB004  Insulin biotech 2020-04-01
```

```
# 22. Sort by the created column and display the first 5 rows
general_information$created <- as.Date(general_information$created)
sorted_df <- general_information[order(general_information$created), ]
head(sorted_df, 5)
```

	drugbank_id	name	type	created
1	DB001	Aspirin	small molecule	2020-01-01
2	DB002	Ibuprofen	small molecule	2020-02-01
3	DB003	Paracetamol	small molecule	2020-03-01
4	DB004	Insulin	biotech	2020-04-01
5	DB005	Morphine	small molecule	2020-05-01

```
# 23. Subset with specific columns and display the first 5 rows
subset_df <- general_information[, c("drugbank_id", "name")]
head(subset_df, 5)
```

	drugbank_id	name
1	DB001	Aspirin
2	DB002	Ibuprofen
3	DB003	Paracetamol
4	DB004	Insulin
5	DB005	Morphine

```
# 24. Merge dataframes and count rows
drug_classification <- db_data$drugs$drug_classification
merged_df <- merge(general_information, drug_classification, by = "drugbank_id")
nrow(merged_df)
```

```
[1] 5
```

```
# 25. Count unique experimental properties (kind)
experimental_properties <- db_data$drugs$experimental_properties
unique_kinds <- unique(experimental_properties$kind)
length(unique_kinds)
```

```
[1] 2
```

```
# 26. Filter for kind "logP" and count rows
logP_df <- experimental_properties[experimental_properties$kind == "logP", ]
nrow(logP_df)
```

```
[1] 5
```

```
# 27. Convert value column to numeric and calculate mean
logP_df$value <- as.numeric(logP_df$value)
mean(logP_df$value, na.rm = TRUE)
```

```
[1] 1.48
```

```
# 28. Calculate summary statistics for logP values
summary(logP_df$value)
```


Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.80	1.20	1.50	1.48	1.80	2.10

```
sd(logP_df$value, na.rm = TRUE)
```

```
[1] 0.5069517
```

```
# 29. Create a histogram of molecular weight values
molecular_weight <- experimental_properties[experimental_properties$kind == "Molecular Weight", ]
molecular_weight$value <- as.numeric(molecular_weight$value)
# clean based on 3 sigma rule
molecular_weight_clean <- molecular_weight[
  abs(molecular_weight$value - mean(molecular_weight$value, na.rm = TRUE)) <= 3 * sd(molecular_weight$value), ]
hist(molecular_weight_clean$value, main = "Histogram of Molecular Weight", xlab = "Molecular Weight")
```

Histogram of Molecular Weight



```
# 30. Filter for kind "Water Solubility" and count unique values
water_solubility_df <- experimental_properties[experimental_properties$kind == "Water Solubility"]
length(unique(water_solubility_df$value))
```

```
[1] 0
```


Data visualization (mainly: ggplot2)

Box plots

```
library(ggplot2)
data(iris)
boxplot(iris$Sepal.Length ~ iris$Species)
```




```
ggplot(mpg, aes(x=class, y=hwy)) +
  geom_boxplot() +
  scale_y_continuous(breaks = seq(10, 45, by=5)) #---diy scale in y-axis
```




```
# ggplot(mpg, aes(x=class, y=hwy)) + geom_boxplot() +scale_y_continuous(breaks = seq(10, 45,
```

Histogram plots


```
hist(iris$Sepal.Length)
```



```
plot(Petal.Length ~ Sepal.Length, data=iris)  
abline(lm(Petal.Length ~ Sepal.Length, data=iris), col="red")
```



```
data(mpg, package='ggplot2')  
ggplot(data=mpg, aes(x=class)) +  
  geom_bar()
```




```
# By default, the geom_bar() just counts the number of cases and displays how many observations
table(mpg$class)
```

2seater	compact	midsize	minivan	pickup	subcompact	suv
5	47	41	11	33	35	62


```
df <- as.data.frame(table(mpg$class))
df
```

	Var1	Freq
1	2seater	5
2	compact	47
3	midsize	41
4	minivan	11
5	pickup	33
6	subcompact	35
7	suv	62

```
ggplot(df, aes(Var1, Freq)) +
  geom_col()
```

```
ggplot(mpg, aes(x=hwy)) + geom_histogram(binwidth = 2)
```



Density plots


```
p1 <- ggplot(mpg, aes(x=hwy, y=after_stat(density))) +  
  geom_histogram(bins=8, fill="blue", alpha=0.5) +  
  labs(title="Histogram of Highway MPG density")  
p2 <- ggplot(mpg, aes(x=hwy)) +  
  geom_density(fill='red', alpha=0.5) +  
  labs(title="Density Plot of Highway MPG")  
library(gridExtra)
```

Attaching package: 'gridExtra'

The following object is masked from 'package:dplyr':


combine

```
grid.arrange(p1, p2, ncol = 2)
```



- multiple plots in one figure

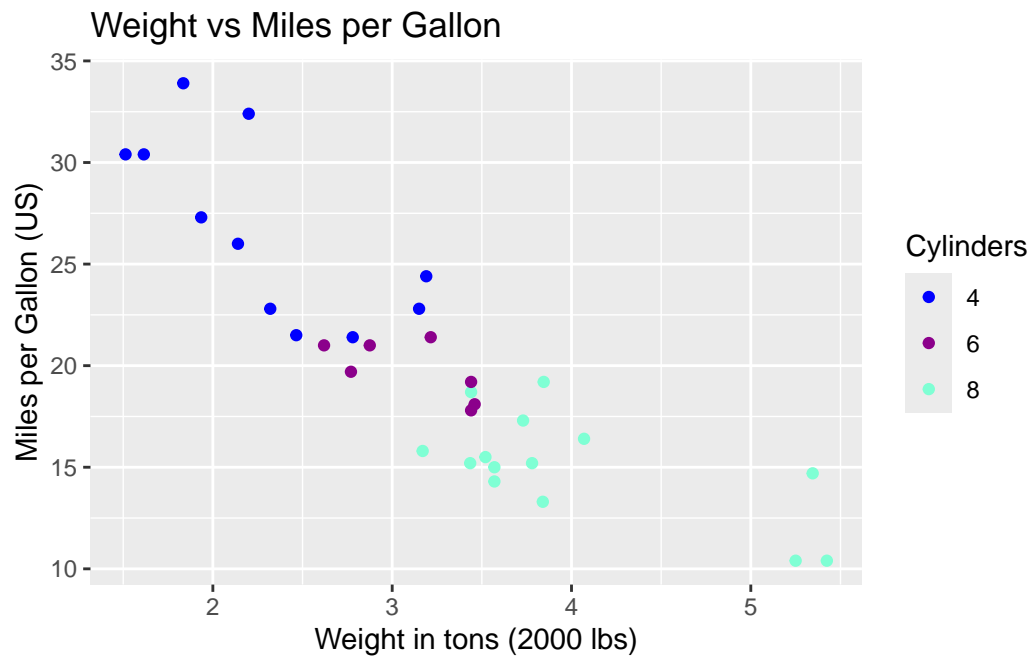
```
ggplot(iris, aes(x = Sepal.Length)) +
  geom_density(aes(fill = Species,color=Species), alpha = 0.5) +
  labs(title = "Density plot") +
  labs(x = "Sepal Length", y = "Density") +
  labs(fill = "area",color="line") # fill is the area,color is the line or dot
```



Scatter plots

```
mtcars$cyl <- factor(mtcars$cyl)


ggplot(mtcars, aes(x=wt, y=mpg, col=cyl)) + geom_point() +
  labs(title='Weight vs Miles per Gallon') +
  labs(x="Weight in tons (2000 lbs)", y="Miles per Gallon (US)" ) +
  labs(color="Cylinders") + # color is dot or line
  scale_color_manual(values=c('blue', 'darkmagenta', 'aquamarine')) # diy color
```



Scatter plots with regression line


```
ggplot(data=iris, aes(x=Sepal.Length, y=Petal.Length, color=Species))+
  geom_point()+# Anything set inside an aes() command will be of the form attribute=Column_Name
  geom_smooth(method="lm") #By default, geom_smooth(method="lm") fits a linear regression line
```

`geom_smooth()` using formula = 'y ~ x'



```
ggplot(data=iris, aes(x=Sepal.Length, y=Petal.Length)) +  
  geom_point(aes(color=Species, shape=Species))+  
  geom_smooth(method="lm")
```

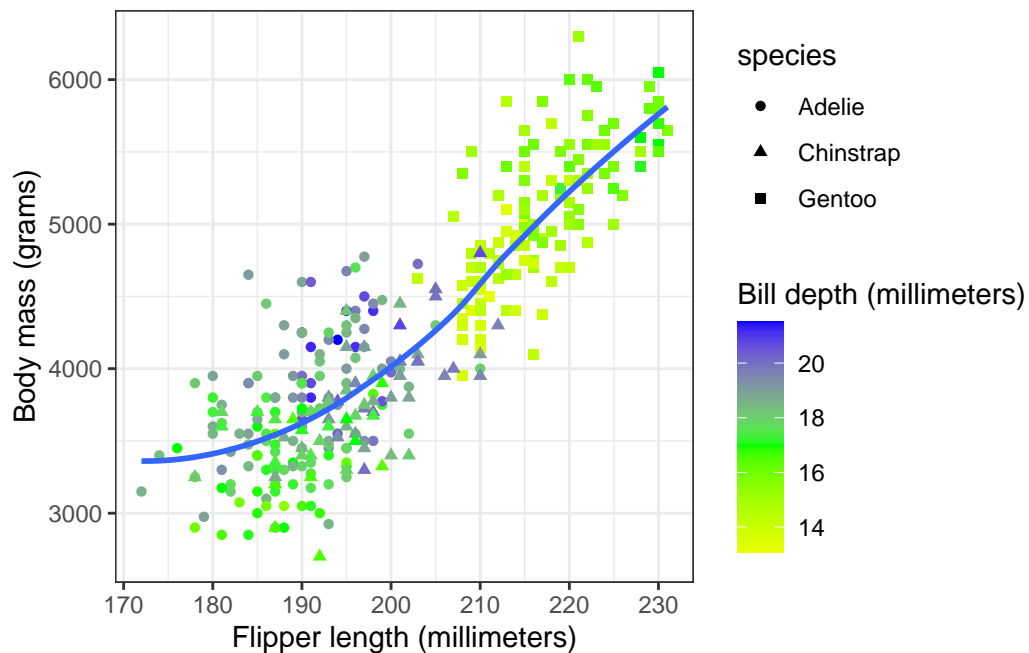
`geom_smooth()` using formula = 'y ~ x'



```
# Zooming in/out--Danger! This removes the data points first!
#ggplot(trees, aes(x=Girth, y=Volume)) +
#  geom_point() +
#  geom_smooth(method='lm') +
#  xlim( 8, 19 ) + ylim(0, 60)
```

```
library(palmerpenguins)
ggplot(penguins, aes(x=flipper_length_mm, y=body_mass_g)) + geom_point(aes(color=bill_depth_mm))
```

```
`geom_smooth()` using method = 'loess' and formula = 'y ~ x'
```



Heat map

```
# data
mine.table <- data.frame(
  Sample.name = rep(paste0("Sample", 1:5), each = 3),
  Class = rep(c("Class1", "Class2", "Class3"), times = 5),
  Abundance = c(0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0, 1.1, 1.2, 1.3, 1.4, 1.5),
  Depth = c(0.5,0.5,0.5, 0.6,0.6,0.6, 0.7,0.7,0.7, 0.8,0.8,0.8, 0.9,0.9,0.9)
)
print(mine.table)
```

	Sample.name	Class	Abundance	Depth
1	Sample1	Class1	0.1	0.5
2	Sample1	Class2	0.2	0.5
3	Sample1	Class3	0.3	0.5
4	Sample2	Class1	0.4	0.6
5	Sample2	Class2	0.5	0.6
6	Sample2	Class3	0.6	0.6
7	Sample3	Class1	0.7	0.7
8	Sample3	Class2	0.8	0.7
9	Sample3	Class3	0.9	0.7

```

10     Sample4 Class1      1.0  0.8
11     Sample4 Class2      1.1  0.8
12     Sample4 Class3      1.2  0.8
13     Sample5 Class1      1.3  0.9
14     Sample5 Class2      1.4  0.9
15     Sample5 Class3      1.5  0.9


```

```

mine.heatmap <- ggplot(data = mine.table, mapping = aes(x = Sample.name, y = Class, fill = Abundance)) +
  geom_tile() + # create the heatmap with tiles+
  scale_y_discrete(limits = rev(levels(factor(mine.table$Class)))) + # reverse the order of
  scale_fill_gradient(low = "white", high = "blue") + # color
  theme_minimal() + # control the theme of the plot
  theme(axis.text.x = element_text(angle = 45, hjust = 1)) + # Rotate the X-axis label for
  labs(x = "Sample Name", # x-axis label
       y = "Class", # y-axis label
       fill = "Abundance")+ # fill legend label by "Abundance"
  ggtitle("Heatmap of Abundance by Sample and Class")+ # add title
  geom_text(aes(label = round(Abundance, 2)), color = "black", size = 3) # add text label

print(mine.heatmap)

```




Create heat map using `facet_grid` to show the data in different panels by depth


```

mine.heatmap <- ggplot(data = mine.table, mapping = aes(x = Sample.name, y = Class, fill = Abundance)) +
  geom_tile() + # create the heatmap with tiles+
  facet_grid(~ Depth, switch = 'x', scales='free', space='free')+ # facet_grid to show the data by depth
  scale_y_discrete(limits = rev(levels(factor(mine.table$Class)))) + # reverse the order of classes
  scale_fill_gradient(low="#FFFFFF", high="#012345")+ # color gradient
  theme_minimal() + # control the theme of the plot
  theme(axis.text.x = element_text(angle = 45, hjust = 1)) + # Rotate the X-axis label for better readability
  labs(x = "Sample Name", # x-axis label
       y = "Class", # y-axis label
       fill = "Abundance")+ # fill legend label by "Abundance"
  ggtitle("Heatmap of Abundance by Sample and Class")+ # add title
  geom_text(aes(label = round(Abundance, 2)), color = "black", size = 3) # add text labels to each tile


print(mine.heatmap)

```



Faceting (make many panels of graphics where each panel represents the same relationship between variables, but something changes between each pane)

```
ggplot(iris, aes(x=Sepal.Length, y=Sepal.Width)) +  
  geom_point() +  
  facet_grid(.~Species) #or facet_grid(Species~.)--Categorical variables of species will be v
```



- Another example

```
library(reshape)
```

Attaching package: 'reshape'

The following object is masked from 'package:dplyr':

rename


The following object is masked from 'package:Matrix':

expand


```
data(tips, package='reshape')
head(tips, 3)
```

	total_bill	tip	sex	smoker	day	time	size
1	16.99	1.01	Female	No	Sun	Dinner	2
2	10.34	1.66	Male	No	Sun	Dinner	3
3	21.01	3.50	Male	No	Sun	Dinner	3

```
ggplot(tips, aes(x = total_bill, y = tip / total_bill)) +
  geom_point() +
  facet_grid( smoker ~ day )
```



```
# 'free_y' means the scale of different panels are adjusted by themselves
ggplot(tips, aes(x = total_bill, y = tip / total_bill)) +
  geom_point() +
  facet_wrap( ~ day, scales='free_y')
```



```
# log scales ---a wrapper of scale_y_continuous() function , trans_new() function
# ggplot(ACS, aes(x=Age, y=Income)) + geom_point() +
# scale_y_log10(breaks=c(1, 10, 100),
#               minor=c(1:10,
#                       seq(10, 100, by=10 ),
#                       seq(100, 1000, by=100))) +
# ylab('Income (1000s of dollars)')
```

- Multi-plot

```
p1 <- ggplot(ChickWeight, aes(x=Time, y=weight, colour=Diet, group=Chick)) +
  geom_line() +
  ggtitle("Growth curve for individual chicks")
# Second plot
p2 <- ggplot(ChickWeight, aes(x=Time, y=weight, colour=Diet)) +
  geom_point(alpha=.3) +
  geom_smooth(alpha=.2, linewidth=1) +
  ggtitle("Fitted growth curve per diet")
# Third plot
p3 <- ggplot(subset(ChickWeight, Time==21), aes(x=weight, colour=Diet)) +
  geom_density() +
  ggtitle("Final weight, by diet")
```

```
# to realize:
# plot1 plot2 plot2
# plot1 plot2 plot2
# plot1 plot3 plot3

my.layout = cbind( c(1,1,1), c(2,2,3), c(2,2,3) ) # each c represents a column in a matrix and
library(Rmisc)
```

Loading required package: plyr

You have loaded plyr after dplyr - this is likely to cause problems.
 If you need functions from both plyr and dplyr, please load plyr first, then dplyr:
 library(plyr); library(dplyr)

Attaching package: 'plyr'

The following objects are masked from 'package:reshape':

rename, round_any

The following object is masked from 'package:mosaic':


count

The following objects are masked from 'package:dplyr':

arrange, count, desc, failwith, id, mutate, rename, summarise,
 summarize

```
Rmisc::multiplot( p1, p2, p3, layout=my.layout)
```

```
`geom_smooth()` using method = 'loess' and formula = 'y ~ x'
```



```
# OR library(ggpubr) https://rpkgs.datanovia.com/ggpubr/.
```

```
library(ggpubr)
```

Attaching package: 'ggpubr'

The following object is masked from 'package:plyr':

mutate

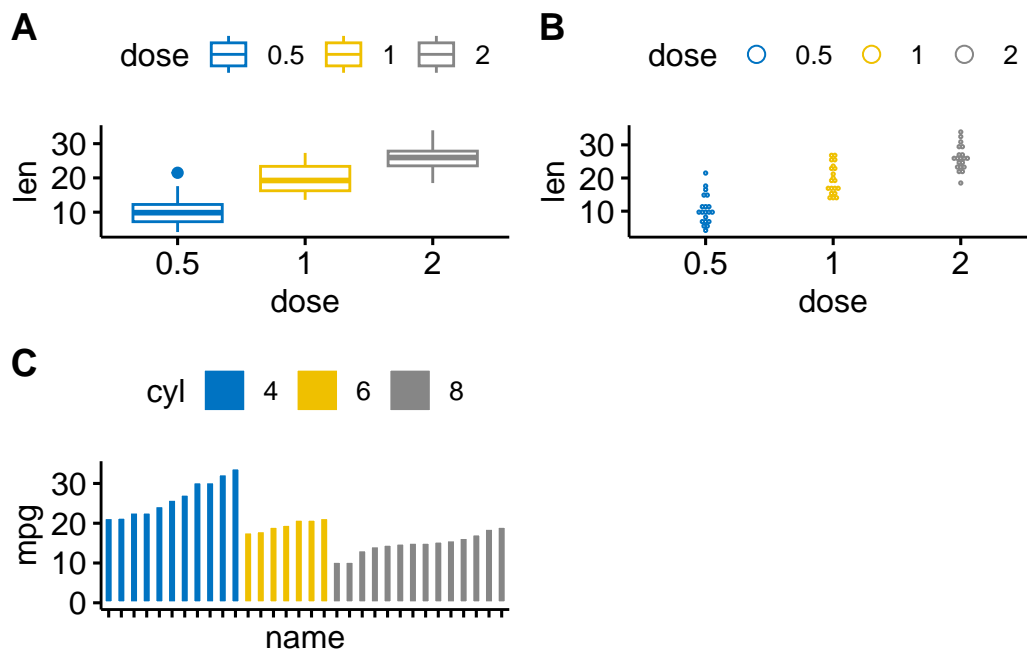
```
# Box plot (bp)
bxp <- ggboxplot(ToothGrowth, x = "dose", y = "len",
  color = "dose", palette = "jco")
# Dot plot (dp)
dp <- ggdotplot(ToothGrowth, x = "dose", y = "len",
  color = "dose", palette = "jco", binwidth = 1)
mtcars$name <- rownames(mtcars)
mtcars$cyl <- as.factor(mtcars$cyl)
bp <- ggbarplot(mtcars, x = "name", y = "mpg",
  fill = "cyl", # change fill color by cyl
  color = "white", # Set bar border colors to white
```

```

    palette = "jco",           # jco journal color palett. see ?ggpar
    sort.val = "asc",          # Sort the value in ascending order
    sort.by.groups = TRUE,     # Sort inside each group
    x.text.angle = 90          # Rotate vertically x axis texts
  ) + font("x.text", size = 8)
# Scatter plots (sp)
sp <- ggscatter(mtcars, x = "wt", y = "mpg",
  add = "reg.line",           # Add regression line
  conf.int = TRUE,            # Add confidence interval
  color = "cyl", palette = "jco", # Color by groups "cyl"
  shape = "cyl"               # Change point shape by groups "cyl"
) +
  stat_cor(aes(color = cyl), label.x = 3) # Add correlation coefficient

ggarrange(bxp, dp, bp + rremove("x.text"),
  labels = c("A", "B", "C"),
  ncol = 2, nrow = 2)

```



```

# Themes
# Rmisc::multiplot( p1 + theme_bw(),           # Black and white
#                   p1 + theme_minimal(),
#                   p1 + theme_dark(),
#                   p1 + theme_light(),

```

```
#                                cols=2 )

#ggsave('p1.png', width=6, height=3, dpi=350)
```

Data manipulation

```
library(dplyr)
library(tidyverse)
```

```
-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
v forcats    1.0.0      v stringr    1.5.1
v lubridate  1.9.4      v tibble     3.2.1
v purrr      1.0.2      v tidyr      1.3.1
v readr      2.1.5

-- Conflicts ----- tidyverse_conflicts() --
x plyr::arrange()      masks dplyr::arrange()
x gridExtra::combine() masks dplyr::combine()
x purrr::compact()     masks plyr::compact()
x plyr::count()        masks mosaic::count(), dplyr::count()
x purrr::cross()       masks mosaic::cross()
x plyr::desc()         masks dplyr::desc()
x mosaic::do()         masks dplyr::do()
x tidyr::expand()      masks reshape::expand(), Matrix::expand()
x plyr::failwith()     masks dplyr::failwith()
x dplyr::filter()      masks stats::filter()
x plyr::id()           masks dplyr::id()
x dplyr::lag()         masks stats::lag()
x ggpubr::mutate()     masks plyr::mutate(), dplyr::mutate()
x tidyr::pack()        masks Matrix::pack()
x dplyr::recode()      masks car::recode()
x plyr::rename()       masks reshape::rename(), dplyr::rename()
x purrr::some()        masks car::some()
x lubridate::stamp()   masks reshape::stamp()
x mosaic::stat()       masks ggplot2::stat()
x plyr::summarise()    masks dplyr::summarise()
x plyr::summarize()    masks dplyr::summarize()
x mosaic::tally()      masks dplyr::tally()
x tidyr::unpack()      masks Matrix::unpack()
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become
```


apply

```
# apply
# Summarize each column by calculating the mean.
apply(iris[,-5],          # what object do we want to apply the function to
      MARGIN=1,          # rows = 1, columns = 2, (same order as [rows, cols])
      FUN=mean           # what function do we want to apply
) %>% head(10)
```

```
[1] 2.550 2.375 2.350 2.350 2.550 2.850 2.425 2.525 2.225 2.400
```

```
average <- apply(
  iris[,-5],          # what object do we want to apply the function to
  MARGIN=2,          # rows = 1, columns = 2, (same order as [rows, cols])
  FUN=mean           # what function do we want to apply
)
iris <- rbind(iris[,-5], average)
iris %>% head(3)
```

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width
1	5.1	3.5	1.4	0.2
2	4.9	3.0	1.4	0.2
3	4.7	3.2	1.3	0.2

There are several variants of the `apply()` function, and the most frequently used ones are `lapply()` and `sapply()`. These two functions apply a given function to each element of a list or vector and returns a corresponding list or vector of results.

```
#lapply
x <- list(a = 1:10, beta = exp(-3:3), logic = c(TRUE,FALSE,FALSE,TRUE))
x
```

\$a

```
[1] 1 2 3 4 5 6 7 8 9 10
```

\$beta

```
[1] 0.04978707 0.13533528 0.36787944 1.00000000 2.71828183 7.38905610
[7] 20.08553692
```

\$logic

```
[1] TRUE FALSE FALSE TRUE
```

```
lapply(x, quantile, probs = 1:3/4) # list
```

```
$a
 25% 50% 75%
3.25 5.50 7.75

$beta
      25%      50%      75%
0.2516074 1.0000000 5.0536690

$logic
 25% 50% 75%
0.0 0.5 1.0
```

```
sapply(x, quantile, probs = 1:3/4) # matrix
```

```
      a      beta logic
25% 3.25 0.2516074  0.0
50% 5.50 1.0000000  0.5
75% 7.75 5.0536690  1.0
```

Tibbles

A tibble, or `tbl_df`, is a modern reimagining of the `data.frame`, keeping what time has proven to be effective, and throwing out what is not. Tibbles are `data.frames` that are lazy and surly: they do less (i.e. they don't change variable names or types, and don't do partial matching) and complain more (e.g. when a variable does not exist). This forces you to confront problems earlier, typically leading to cleaner, more expressive code. Tibbles also have an enhanced `print()` method which makes them easier to use with large datasets containing complex objects.

```
data <- data.frame(a = 1:3, b = letters[1:3], c = Sys.Date() - 1:3)
data
```

```
  a b      c
1 1 a 2025-06-07
2 2 b 2025-06-06
3 3 c 2025-06-05
```

```
as_tibble(data)
```

```
# A tibble: 3 x 3
      a b      c
  <int> <chr> <date>
1     1 a 2025-06-07
2     2 b 2025-06-06
3     3 c 2025-06-05
```

%>%

The pipe operator %>% is used to pass the result of one function to the next function in a chain, making the code more readable and concise. For example, if we wanted to start with x, and first apply function f(), then g(), and then h(), the usual R command would be h(g(f(x))) which is hard to read because you have to start reading at the innermost set of parentheses. Using the pipe command %>%, this sequence of operations becomes x %>% f() %>% g() %>% h().

select

```
# Correct usage of select() within a pipeline
starwars %>% select(-ends_with('color'))
```

filter

```
library(dplyr)

# Filter rows where species is "Droid" and mass is greater than or equal to 100
filtered_data <- starwars %>% filter(species == "Droid", mass < 100)
print(filtered_data)
```

```
# A tibble: 3 x 14
  name height mass hair_color skin_color eye_color birth_year sex gender
  <chr>  <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
1 C-3PO   167    75 <NA>      gold        yellow        112 none masculine
2 R2-D2    96    32 <NA>      white, blue red          33 none masculine
```

```
3 R5-D4      97      32 <NA>      white, red  red      NA none  masculine
# i 5 more variables: homeworld <chr>, species <chr>, films <list>,
#   vehicles <list>, starships <list>
```

slice

This function is used to select rows by their position in the data frame. It can be used to select specific rows or a range of rows.

filter rows based on row number:

```
starwars %>% slice(2:4)
```

```
# A tibble: 3 x 14
  name      height  mass hair_color skin_color eye_color birth_year sex  gender
  <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
1 C-3PO      167    75 <NA>      gold        yellow        112  none  mascu~
2 R2-D2       96    32 <NA>      white, bl~ red          33  none  mascu~
3 Darth Va~   202   136 none      white       yellow        41.9 male  mascu~
# i 5 more variables: homeworld <chr>, species <chr>, films <list>,
#   vehicles <list>, starships <list>
```

arrange

This function is used to sort the rows of a data frame by one or more columns. The default sorting of the number in the dataset is in ascending order, but you can use the `desc()` function to sort in descending order.

```
starwars %>% arrange(desc(name)) #The default sorting is in ascending order
```

```
# A tibble: 87 x 14
  name      height  mass hair_color skin_color eye_color birth_year sex  gender
  <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
1 Zam Wes~   168    55 blonde    fair, gre~ yellow        NA fema~ femin~
2 Yoda        66    17 white     green      brown      896 male  mascu~
3 Yarael ~   264    NA none      white     yellow        NA male  mascu~
4 Wilhuff~   180    NA auburn, g~ fair       blue        64 male  mascu~
5 Wicket ~    88    20 brown     brown     brown         8 male  mascu~
6 Wedge A~   170    77 brown     fair      hazel        21 male  mascu~
7 Watto     137    NA black     blue, gre~ yellow        NA male  mascu~
```

```

 8 Wat Tam~      193      48 none      green, gr~ unknown      NA male  mascu~
 9 Tion Me~      206      80 none      grey      black      NA male  mascu~
10 Taun We      213      NA none      grey      black      NA fema~ femin~
# i 77 more rows
# i 5 more variables: homeworld <chr>, species <chr>, films <list>,
#   vehicles <list>, starships <list>

```

```
starwars %>% arrange(desc(height)) %>% head(3)
```

```

# A tibble: 3 x 14
  name      height  mass hair_color skin_color eye_color birth_year sex  gender
  <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
1 Yarael P~    264     NA none      white      yellow      NA male  mascu~
2 Tarfful      234    136 brown      brown      blue      NA male  mascu~
3 Lama Su      229     88 none      grey      black      NA male  mascu~
# i 5 more variables: homeworld <chr>, species <chr>, films <list>,
#   vehicles <list>, starships <list>

```

```

dd <- data.frame(
  Trt = factor(c("High", "Med", "High", "Low"),
              levels = c("Low", "Med", "High")), # level
  y = c(8, 3, 9, 9),
  z = c(1, 1, 1, 2))
dd %>% arrange(Trt, desc(y))

```

```

      Trt y z
1 Low  9 2
2 Med  3 1
3 High 9 1
4 High 8 1

```

mutate

This function is used to create new columns or modify existing columns in a data frame. It allows you to perform calculations and transformations on the data.

```

# select using the old columns
starwars$bmi = starwars$mass / ((starwars$height / 100) ^ 2)
starwars %>% select(name, bmi) %>% head(3)

```

```

# mutate avoids all the starwars$

starwars$bmi <- NULL
starwars %>%
  mutate(bmi = mass / ((height / 100) ^ 2)) %>%
  select(name, bmi) %>% head(3)

# mutate_at() and mutate_if() allow us to apply a function to a particular column and save the result

subset <- starwars %>%
  mutate(square_height = (height / 100) ^ 2,
         bmi = mass / square_height) %>%
  select(name, square_height, bmi)
subset %>% head(3)

subset %>% mutate_if(is.numeric, round, digits=0) # here, is.numeric is the condition

subset %>% mutate_at(2:3, round, digits=0) %>% head() # column 2 3

# Apply the transformation to columns 2 and 3 for rows 1 to 3
result <- subset %>%
  mutate_at(2:3, ~ifelse(row_number() %in% 1:3, round(., digits = 0), .))

subset %>% mutate(avg.example = select(., square_height:bmi) %>% rowMeans())

```

summarise

This function is used to create a summary table. It reduces the data frame to a single row containing summary statistics.

```

starwars %>% summarise(mean.height=mean(height, na.rm=T), sd.height=sd(height, na.rm=T))

# apply the same statistic to each column
starwars %>% select(height:mass) %>% summarise_all(list(min=min, max=max), na.rm=T)
starwars %>% summarise_if(is.numeric, list(min=min, max=max), na.rm = T)

```

group_by

This function is used to group the data frame by one or more columns. It is often used in combination with `summarise()` to calculate summary statistics for each group.

```

library(dplyr)
library(palmerpenguins)
table(penguins$sex, penguins$species)
penguins %>%
  filter(!is.na(sex)) %>%
  group_by(sex, species) %>%
  summarise(n = n(),
            mean.flipper = mean(flipper_length_mm),
            sd.flipper = sd(flipper_length_mm),
            .groups='keep') %>%
  head(3)

```

examples

Find the flight with the longest departure delay among flights from the same origin and destination (use `filter()`). Relocate the origin, destination, and departure delay to the first three columns and sort by origin and dest.

```

flights %>%
  filter(!is.na(dep_delay)) %>%
  group_by(origin, dest) %>%
  filter(dep_delay == max(dep_delay)) %>%
  relocate(origin, dest, dep_delay) %>%
  arrange(origin, dest)

```

Find the flight with the longest departure delay among flights from the same origin and destination (use `top_n()` or `slice_max()`). Relocate the origin, destination, and departure delay to the first three columns and sort by origin and dest.

```

flights %>%
  filter(!is.na(dep_delay)) %>%
  group_by(origin, dest) %>%
  top_n(1, dep_delay) %>% # or using slice_max(dep_delay) %>%
  relocate(origin, dest, dep_delay) %>%
  arrange(origin, dest)

```

How do departure delays vary at different times of the day? Summarize the averaged departure delays by hours and create a new column named as `dep_delay_level` which `cut()` the averaged departure delays into three levels (low, median, and high).

```
flights %>%
  group_by(hour) %>%
  summarize(avg_dep_delay = mean(dep_delay, na.rm = TRUE)) %>%
  mutate(dep_delay_level = cut(avg_dep_delay, breaks=3, c('low', 'median', 'high')))
```

How do departure delays vary at different times of the day? Illustrate your answer with a `geom_smooth()` plot.

```
flights %>%
  group_by(hour) %>%
  summarize(avg_dep_delay = mean(dep_delay, na.rm = TRUE)) %>%
  ggplot(aes(x = hour, y = avg_dep_delay)) + geom_smooth()
```

```
# ways to delete the blanks
students %>%
  rename(
    student_id = `Student ID`,
    full_name = `Full Name`
  ) %>% head(3)
```

```
read_csv(
  "# A comment I want to skip
  x,y,z
  1,2,3",
  comment = "#"
)
```

```
library(tidyr)
grade.book <- rbind(
  data.frame(name='Alison', HW.1=8, HW.2=5, HW.3=8, HW.4=4),
  data.frame(name='Brandon', HW.1=5, HW.2=3, HW.3=6, HW.4=9),
  data.frame(name='Charles', HW.1=9, HW.2=7, HW.3=9, HW.4=10))
grade.book
```

	name	HW.1	HW.2	HW.3	HW.4
1	Alison	8	5	8	4
2	Brandon	5	3	6	9
3	Charles	9	7	9	10


```
tidy.scores <- grade.book %>%
  pivot_longer(
    cols = starts_with("HW"),
    names_to = "Homework",
    values_to = "Score"
  )
tidy.scores
```

```
# A tibble: 12 x 3
  name      Homework Score
  <chr>    <chr>    <dbl>
1 Alison HW.1      8
2 Alison HW.2      5
3 Alison HW.3      8
4 Alison HW.4      4
5 Brandon HW.1     5
6 Brandon HW.2     3
7 Brandon HW.3     6
8 Brandon HW.4     9
9 Charles HW.1     9
10 Charles HW.2     7
11 Charles HW.3     9
12 Charles HW.4    10
```

```
tidy.scores %>% pivot_wider(names_from=Homework, values_from=Score)
```

```
# A tibble: 3 x 5
  name      HW.1 HW.2 HW.3 HW.4
  <chr>    <dbl> <dbl> <dbl> <dbl>
1 Alison      8     5     8     4
2 Brandon     5     3     6     9
3 Charles     9     7     9    10
```

```
# table joins
```

```
Fish.Data <- tibble(
  Lake_ID = c('A','A','B','B','C','C'),
  Fish.Weight=rnorm(6, mean=260, sd=25) ) # make up some data
Fish.Data
```

```
# A tibble: 6 x 2
  Lake_ID Fish.Weight
  <chr>      <dbl>
1 A          273.
2 A          262.
3 B          284.
4 B          259.
5 C          240.
6 C          280.
```

```
Lake.Data <- tibble(
  Lake_ID = c('B','C','D'),
  Lake_Name = c('Lake Elaine', 'Mormon Lake', 'Lake Mary'),
  pH=c(6.5, 6.3, 6.1),
  area = c(40, 210, 240),
  avg_depth = c(8, 10, 38))
Lake.Data
```

```
# A tibble: 3 x 5
  Lake_ID Lake_Name      pH area avg_depth
  <chr>    <chr>      <dbl> <dbl>    <dbl>
1 B      Lake Elaine  6.5    40         8
2 C      Mormon Lake  6.3   210        10
3 D      Lake Mary   6.1   240        38
```

```
full_join(Fish.Data, Lake.Data)
```

Joining with `by = join_by(Lake_ID)`

```
# A tibble: 7 x 6
  Lake_ID Fish.Weight Lake_Name      pH area avg_depth
  <chr>      <dbl> <chr>      <dbl> <dbl>    <dbl>
1 A          273. <NA>      NA     NA     NA
2 A          262. <NA>      NA     NA     NA
3 B          284. Lake Elaine  6.5    40         8
4 B          259. Lake Elaine  6.5    40         8
5 C          240. Mormon Lake  6.3   210        10
6 C          280. Mormon Lake  6.3   210        10
7 D           NA Lake Mary   6.1   240        38
```

```
left_join(Fish.Data, Lake.Data)
```

Joining with `by = join_by(Lake_ID)`

```
# A tibble: 6 x 6
  Lake_ID Fish.Weight Lake_Name      pH area avg_depth
  <chr>      <dbl> <chr>      <dbl> <dbl>    <dbl>
1 A          273. <NA>        NA      NA        NA
2 A          262. <NA>        NA      NA        NA
3 B          284. Lake Elaine  6.5    40         8
4 B          259. Lake Elaine  6.5    40         8
5 C          240. Mormon Lake  6.3   210        10
6 C          280. Mormon Lake  6.3   210        10
```

```
inner_join(Fish.Data, Lake.Data)
```

Joining with `by = join_by(Lake_ID)`

```
# A tibble: 4 x 6
  Lake_ID Fish.Weight Lake_Name      pH area avg_depth
  <chr>      <dbl> <chr>      <dbl> <dbl>    <dbl>
1 B          284. Lake Elaine  6.5    40         8
2 B          259. Lake Elaine  6.5    40         8
3 C          240. Mormon Lake  6.3   210        10
4 C          280. Mormon Lake  6.3   210        10
```

```
mutate(
  week = parse_number(week)
)
```

how many data points are in the data set

```
gender_year <- Survey %>%
  filter(!is.na(Year)) %>%
  group_by(Sex, Year) %>%
  count() %>%
  rename(nu=n)
gender_year
gender_year %>% pivot_wider(names_from = Year, values_from = nu)
```

```

who2 %>%
  head(3)
who2 <- who2 %>%
  pivot_longer(
    cols = !(country:year),
    names_to = c("diagnosis", "gender", "age"),
    names_sep = "_",
    values_to = "count"
  ) %>%
  filter(!is.na(count))
who2

```

```

left_join(feb14_VX, airports, by=c('dest'='faa'))

```

```

library(psych)
drug_prop <- drug_prop %>%
  filter(class == 'Carboxylic acids and derivatives')
drug_prop %>%
  select(logP, logS, water_solubility) %>%
  pairs.panels()

```

ps:the comparison between whether to use %>% or not

```

# %>%
penguins %>%
  filter(!is.na(sex)) %>%
  group_by(sex, species) %>%
  mutate(Sum.Sq.Cells = (flipper_length_mm - mean(flipper_length_mm))^2) %>%
  select(sex, species, flipper_length_mm, Sum.Sq.Cells) %>% head()

# not use %>%
head(
  select(mutate(group_by(filter(penguins, !is.na(sex)), sex, species),
    Sum.Sq.Cells = (flipper_length_mm - mean(flipper_length_mm))^2),
    sex, species, flipper_length_mm, Sum.Sq.Cells))

```

```

library(nycflights13)
str(nycflights13::flights)
# the order of group_by and summarize matters

```

```
flights %>%
  group_by(carrier) %>%
  summarize(avg_dep_delay = mean(dep_delay, na.rm = TRUE)) %>%
  arrange(desc(avg_dep_delay))
```

Control flow

while loop

```
x <- 1
while (x < 10) {
  print(x)
  x <- x + 1
}
```

for loop

Fibonacci sequence

```
F <- rep(0, 10)
F[1] <- 0
F[2] <- 1
cat('F = ', F, '\n')
```

```
F = 0 1 0 0 0 0 0 0 0 0
```

```
for( n in 3:10 ){
  F[n] <- F[n-1] + F[n-2]
  cat('F = ', F, '\n')
}
```

```
F = 0 1 1 0 0 0 0 0 0 0
F = 0 1 1 2 0 0 0 0 0 0
F = 0 1 1 2 3 0 0 0 0 0
F = 0 1 1 2 3 5 0 0 0 0
F = 0 1 1 2 3 5 8 0 0 0
```

```
F = 0 1 1 2 3 5 8 13 0 0
F = 0 1 1 2 3 5 8 13 21 0
F = 0 1 1 2 3 5 8 13 21 34
```


bootstrap estimate of a sampling distribution

- bootstrap estimate of a sampling distribution

```
library(dplyr)
library(ggplot2)
SampDist <- data.frame()

for (i in 1:1000){
  SampDist <- trees %>%
    slice_sample(n=30, replace =TRUE) %>%
    dplyr::summarise(xbar=mean(Height)) %>%
    rbind(SampDist)
}
ggplot(SampDist,aes(x=xbar)) +
  geom_histogram()
```

`stat_bin()` using `bins = 30`. Pick better value with `binwidth`.



Functions

Functions construction

copy of x

```
k <- 3
example.func <- function(x){
  x <- sort(x)
  if (k > 1){
    print(x)
  }
}
x <- c(3,1,5,4,2)
example.func(x)
```

```
[1] 1 2 3 4 5
```

```
x # x is changed inside the function but not outside the function
```

```
[1] 3 1 5 4 2
```

Ellipses

```
# a function that draws the regression line and confidence interval
# notice it doesn't return anything, all it does is draw a plot
show.lm <- function(m, interval.type='confidence', fill.col='light grey', ...){
  x <- m$model[,2]      # extract the predictor variable
  y <- m$model[,1]      # extract the response
  pred <- predict(m, interval=interval.type)
  plot(x, y, ...)
  polygon( c(x,rev(x)), # draw the ribbon defined
           c(pred[, 'lwr'], rev(pred[, 'upr'])), # by lwr and upr - polygon
           col='light grey') # fills in the region defined by
  lines(x, pred[, 'fit']) # a set of vertices, need to reverse
  points(x, y) # the uppers to make a nice figure
}
```

Linear regression and multiple linear regression (Lab 10)

1. Load the `bloodpress.txt`

```
bloodpress <- read.table("bloodpress.txt", header=T)
bloodpress
```

	Pt	BP	Age	Weight	BSA	Dur	Pulse	Stress
1	1	105	47	85.4	1.75	5.1	63	33
2	2	115	49	94.2	2.10	3.8	70	14
3	3	116	49	95.3	1.98	8.2	72	10
4	4	117	50	94.7	2.01	5.8	73	99
5	5	112	51	89.4	1.89	7.0	72	95
6	6	121	48	99.5	2.25	9.3	71	10
7	7	121	49	99.8	2.25	2.5	69	42
8	8	110	47	90.9	1.90	6.2	66	8
9	9	110	49	89.2	1.83	7.1	69	62
10	10	114	48	92.7	2.07	5.6	64	35
11	11	114	47	94.4	2.07	5.3	74	90
12	12	115	49	94.1	1.98	5.6	71	21
13	13	114	50	91.6	2.05	10.2	68	47
14	14	106	45	87.1	1.92	5.6	67	80
15	15	125	52	101.3	2.19	10.0	76	98
16	16	114	46	94.5	1.98	7.4	69	95
17	17	106	46	87.0	1.87	3.6	62	18
18	18	113	46	94.5	1.90	4.3	70	12
19	19	110	48	90.5	1.88	9.0	71	99
20	20	122	56	95.7	2.09	7.0	75	99

2. Use `pairs.panels()` function from `psych` package to draw scatterplots, histograms, and calculate correlations between variables.

```
library(psych)
```

Attaching package: 'psych'

The following objects are masked from 'package:mosaic':

`logit`, `rescale`

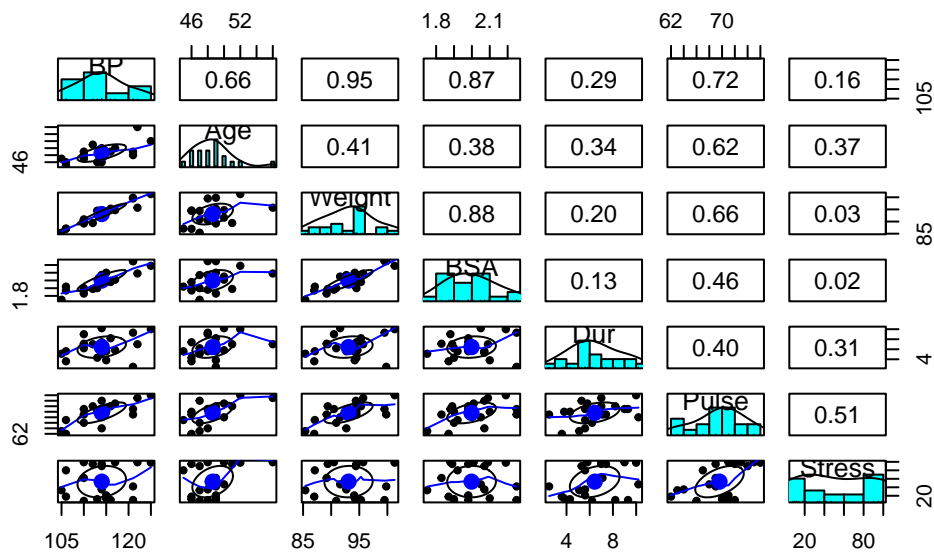
The following object is masked from 'package:car':

logit

The following objects are masked from 'package:ggplot2':

%+%, alpha

```
pairs.panels(bloodpress[, -1])
```



3. Fit a simple linear regression model of BP vs Stress. Is Stress significant?

```
model.1 <- lm(BP ~ Stress, data=bloodpress)
summary(model.1)
```

Call:

```
lm(formula = BP ~ Stress, data = bloodpress)
```

Residuals:

Min	1Q	Median	3Q	Max
-8.6394	-3.3014	0.0722	2.2181	9.9287

Coefficients:

Estimate	Std. Error	t value	Pr(> t)
----------	------------	---------	----------

```
(Intercept) 112.71997    2.19345   51.389   <2e-16 ***
Stress       0.02399    0.03404    0.705     0.49
```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 5.502 on 18 degrees of freedom

Multiple R-squared: 0.02686, Adjusted R-squared: -0.0272

F-statistic: 0.4969 on 1 and 18 DF, p-value: 0.4899

4. Fit a simple linear regression model of BP vs Weight.

```
model.2 <- lm(BP ~ Weight, data=bloodpress)
summary(model.2)
```

Call:

```
lm(formula = BP ~ Weight, data = bloodpress)
```

Residuals:

```
      Min       1Q   Median       3Q      Max
-2.6933 -0.9318 -0.4935  0.7703  4.8656
```

Coefficients:

```
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  2.20531     8.66333   0.255    0.802
Weight       1.20093     0.09297  12.917 1.53e-10 ***
```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.74 on 18 degrees of freedom

Multiple R-squared: 0.9026, Adjusted R-squared: 0.8972

F-statistic: 166.9 on 1 and 18 DF, p-value: 1.528e-10

5. Fit a simple linear regression model of BP vs BSA.

```
model.3 <- lm(BP ~ BSA, data=bloodpress)
summary(model.3)
```

Call:

```
lm(formula = BP ~ BSA, data = bloodpress)
```

Residuals:

Min	1Q	Median	3Q	Max
-5.314	-1.963	-0.197	1.934	4.831

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	45.183	9.392	4.811	0.00014 ***
BSA	34.443	4.690	7.343	8.11e-07 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.79 on 18 degrees of freedom

Multiple R-squared: 0.7497, Adjusted R-squared: 0.7358

F-statistic: 53.93 on 1 and 18 DF, p-value: 8.114e-07

6. Fit a multiple linear regression model of BP vs Weight + BSA. Is BSA still significant? Why?

```
model.4 <- lm(BP ~ Weight + BSA, data=bloodpress)
summary(model.4)
```

Call:

```
lm(formula = BP ~ Weight + BSA, data = bloodpress)
```

Residuals:

Min	1Q	Median	3Q	Max
-1.8932	-1.1961	-0.4061	1.0764	4.7524

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	5.6534	9.3925	0.602	0.555
Weight	1.0387	0.1927	5.392	4.87e-05 ***
BSA	5.8313	6.0627	0.962	0.350

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.744 on 17 degrees of freedom

Multiple R-squared: 0.9077, Adjusted R-squared: 0.8968

F-statistic: 83.54 on 2 and 17 DF, p-value: 1.607e-09

8. Predict BP for Weight=92 and BSA=2 for the two simple linear regression models and the multiple linear regression model, by hand and by `predict()` function.

```
2.20531 + 1.20093 * 92
```

```
[1] 112.6909
```

```
predict(model.2,  
        newdata=data.frame(Weight=92))
```

```
1  
112.691
```

```
45.183 + 34.443 * 2
```

```
[1] 114.069
```

```
predict(model.3,  
        newdata=data.frame(BSA=2))
```

```
1  
114.0689
```

```
5.6534 + 1.0387 * 92 + 5.8313 * 2
```

```
[1] 112.8764
```

```
predict(model.4,  
        newdata=data.frame(Weight=92, BSA=2))
```

```
1  
112.8794
```

7. Fit a multiple linear regression model of BP vs Age + Weight. Set argument `x` and `y` as `TRUE`. Save the output of `lm()` as `model.5`. How do we interpret each estimated coefficients?

```
model.5 <- lm(BP ~ Age + Weight, data=bloodpress, x=TRUE, y=TRUE)  
summary(model.5)
```

Call:

```
lm(formula = BP ~ Age + Weight, data = bloodpress, x = TRUE,  
   y = TRUE)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.89968	-0.35242	0.06979	0.35528	0.82781

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-16.57937	3.00746	-5.513	3.80e-05 ***
Age	0.70825	0.05351	13.235	2.22e-10 ***
Weight	1.03296	0.03116	33.154	< 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.5327 on 17 degrees of freedom

Multiple R-squared: 0.9914, Adjusted R-squared: 0.9904

F-statistic: 978.2 on 2 and 17 DF, p-value: < 2.2e-16

8. Use the `plot_ly` function in the `plotly` package to create a 3D scatterplot of the data with the fitted plane for a multiple linear regression model of BP vs Age + Weight.

```
library(plotly)  
plot_ly(x=bloodpress$Age, y=bloodpress$Weight, z=bloodpress$BP, type='scatter3d', mode='mark
```

9. Extract the matrix `x` and `y` of `model.5` and assign it to a new object `X` and `y`. Remember, if you save the output of `lm()` as an object, this object contains many elements. After we set `x=TRUE` and `y=TRUE` in question 8, we can find `x` and `y` in this list.

```
X <- model.5$x  
y <- model.5$y
```

10. Calculate $X^T X$, $X^T y$, $(X^T X)^{-1}$, and $(X^T X)^{-1} X^T y$. Use `t()` for transpose, `%*%` for matrix multiplication, and `solve()` for inverse of matrix. For the last one, is your result same as the estimated values you obtained in question 7? –Of course!

```
t(X) %*% X
```

	(Intercept)	Age	Weight
(Intercept)	20.0	972.0	1861.8
Age	972.0	47358.0	90566.6
Weight	1861.8	90566.6	173665.4

```
t(X) %*% y
```

	[,1]
(Intercept)	2280.0
Age	110978.0
Weight	212666.1

```
solve(t(X) %*% X)
```

	(Intercept)	Age	Weight
(Intercept)	31.8748075	-0.267669593	-0.202127676
Age	-0.2676696	0.010092130	-0.002393468
Weight	-0.2021277	-0.002393468	0.003420885

```
solve(t(X) %*% X) %*% (t(X) %*% y)
```

	[,1]
(Intercept)	-16.5793694
Age	0.7082515
Weight	1.0329611

11. Use the `anova` function to display the ANOVA table with sequential (type I) sums of squares for the `model.5`.

$$SS_{\text{Variable}} = \sum_{i=1}^n (\hat{y}_{\text{Variable},i} - \bar{y})^2$$

$\hat{y}_{\text{Variable},i}$ is the model including only variable i .

F is the ratio of the mean square for the variable to the mean square for residuals.

If ($F \approx 1$) : It indicates that the sizes of `MS_Variable` and `MS_Residuals` are approximately the same, suggesting that the explanatory power of the independent variable for the dependent variable is comparable to the random error, and the null hypothesis ((H_0)) cannot be rejected.

If ($F \gg 1$) : It indicates that `MS_Variable` is significantly greater than `MS_Residuals`, suggesting that the independent variable has a significant influence on the dependent variable, and the null hypothesis ((H_0)) can be rejected.

```
anova(model.5)
```

Analysis of Variance Table

Response: BP

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
Age	1	243.266	243.266	857.29	5.481e-16 ***
Weight	1	311.910	311.910	1099.20	< 2.2e-16 ***
Residuals	17	4.824	0.284		

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```
# remark
```

```
sum((model.5$y-mean(model.5$y))^2) == 243.266+311.910+4.824
```

```
[1] TRUE
```

$$SS_{\text{Total}} = SS_{\text{Age}} + SS_{\text{Weight}} + SS_{\text{Residuals}}$$

12. Use the `residuals` element in fitted model or `residuals()` function to extract the fitted residuals. Calculate the sum of square of these residual values. Extract the `df.residual` element in fitted model and use the above elements to calculate the MSE. Is your result same as the `anova()` output?

```
sum((model.5$residuals)^2)/model.5$df.residual
```

```
[1] 0.2837604
```

13. Fit a multiple linear regression model of BP vs Age + Weight + Pulse. Save the output of `lm()` as `model.6`.

```
model.6 <- lm(BP ~ Age + Weight + Pulse, data=bloodpress)
summary(model.6)
```

Call:

```
lm(formula = BP ~ Age + Weight + Pulse, data = bloodpress)
```

Residuals:

Min	1Q	Median	3Q	Max
-----	----	--------	----	-----

```
-0.71174 -0.45422 -0.01909 0.41745 0.88743
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-16.69000	2.93761	-5.681	3.40e-05	***
Age	0.75018	0.06074	12.350	1.36e-09	***
Weight	1.06135	0.03695	28.722	3.40e-15	***
Pulse	-0.06566	0.04852	-1.353	0.195	

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.5201 on 16 degrees of freedom

Multiple R-squared: 0.9923, Adjusted R-squared: 0.9908

F-statistic: 684.7 on 3 and 16 DF, p-value: < 2.2e-16

14. Use `anova()` function to obtain the ANOVA table for `model.6`. We may consider `model.6` as full model, and `model.5` as reduced model in this question. Based on the obtained ANOVA table and the output of question 11, calculate the F-statistic for testing the reduced model by hand. You may use the Residuals Sum sq and the corresponding Residuals Df from both tables. Then, calculate the p-value using `pf()` function, don't forget about the `lower.tail`.

```
anova(model.6)
```

Analysis of Variance Table

Response: BP

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
Age	1	243.266	243.266	899.2446	1.726e-15 ***
Weight	1	311.910	311.910	1152.9909	2.433e-16 ***
Pulse	1	0.496	0.496	1.8319	0.1947
Residuals	16	4.328	0.271		

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```
Fstat <- (4.824-4.328)/(17-16) / (4.328/16)
Fstat
```

```
[1] 1.833641
```



```
pf(Fstat, 1, 16, lower.tail = F)
```

```
[1] 0.1945157
```

15. Use `anova()` function to do the F-test on `model.5` and `model.6`. Compare the output with your answers of question 14. What is the conclusion of the F-test?

$$RSS = \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

Sum of Sq = $RSS(\text{Model 1}) - RSS(\text{Model 2})$ (It represents the variation in the interpretation of the dependent variable by the newly added variable Pulse.)

$$F = \frac{\text{Sum of Sq/Df}}{RSS(\text{Model 2})/\text{Res.Df}(\text{Model 2})} F = \frac{0.49557}{0.270525} \approx 1.8319$$

```
anova(model.5, model.6)
```

Analysis of Variance Table


Model 1: BP ~ Age + Weight

Model 2: BP ~ Age + Weight + Pulse

	Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
1	17	4.8239				
2	16	4.3284	1	0.49557	1.8319	0.1947

16. Plot the qqPlot for residuals of `model.5`. What is the x-axis and y-axis of the qqPlot? What can we say about the qqPlot?


```
library(car)
qqPlot(model.5$residuals)
```



```
[1] 19 10
```


17. Plot the residual vs fitted plot of `model.5`. You may extract `fitted.values` from `model.5` and use it as `x` in `plot()`.



```
plot(x=model.5$fitted.values, y=model.5$residuals)
```




18. Directly use `plot()` function on `model.5`.

```
plot(model.5)
```







19. Load the `hospital_infct.txt` data and select observations with `Stay <= 14`.

```
infectionrisk <- read.table("/Users/luyu/Desktop/NOTESAPH101/hospital_infct.txt", header=T)
infectionrisk <- infectionrisk[infectionrisk$Stay<=14,]
infectionrisk
```

	ID	Stay	Age	InfctRsk	Culture	Xray	Beds	MedSchool	Region	Census	Nurses
1	1	7.13	55.7	4.1	9.0	39.6	279	2	4	207	241
2	2	8.82	58.2	1.6	3.8	51.7	80	2	2	51	52
3	3	8.34	56.9	2.7	8.1	74.0	107	2	3	82	54
4	4	8.95	53.7	5.6	18.9	122.8	147	2	4	53	148
5	5	11.20	56.5	5.7	34.5	88.9	180	2	1	134	151
6	6	9.76	50.9	5.1	21.9	97.0	150	2	2	147	106
7	7	9.68	57.8	4.6	16.7	79.0	186	2	3	151	129
8	8	11.18	45.7	5.4	60.5	85.8	640	1	2	399	360
9	9	8.67	48.2	4.3	24.4	90.8	182	2	3	130	118
10	10	8.84	56.3	6.3	29.6	82.6	85	2	1	59	66
11	11	11.07	53.2	4.9	28.5	122.0	768	1	1	591	656
12	12	8.30	57.2	4.3	6.8	83.8	167	2	3	105	59
13	13	12.78	56.8	7.7	46.0	116.9	322	1	1	252	349
14	14	7.58	56.7	3.7	20.8	88.0	97	2	2	59	79
15	15	9.00	56.3	4.2	14.6	76.4	72	2	3	61	38
16	16	11.08	50.2	5.5	18.6	63.6	387	2	3	326	405
17	17	8.28	48.1	4.5	26.0	101.8	108	2	4	84	73
18	18	11.62	53.9	6.4	25.5	99.2	133	2	1	113	101
19	19	9.06	52.8	4.2	6.9	75.9	134	2	2	103	125

20	20	9.35	53.8	4.1	15.9	80.9	833	2	3	547	519
21	21	7.53	42.0	4.2	23.1	98.9	95	2	4	47	49
22	22	10.24	49.0	4.8	36.3	112.6	195	2	2	163	170
23	23	9.78	52.3	5.0	17.6	95.9	270	1	1	240	198
24	24	9.84	62.2	4.8	12.0	82.3	600	2	3	468	497
25	25	9.20	52.2	4.0	17.5	71.1	298	1	4	244	236
26	26	8.28	49.5	3.9	12.0	113.1	546	1	2	413	436
27	27	9.31	47.2	4.5	30.2	101.3	170	2	1	124	173
28	28	8.19	52.1	3.2	10.8	59.2	176	2	1	156	88
29	29	11.65	54.5	4.4	18.6	96.1	248	2	1	217	189
30	30	9.89	50.5	4.9	17.7	103.6	167	2	2	113	106
31	31	11.03	49.9	5.0	19.7	102.1	318	2	1	270	335
32	32	9.84	53.0	5.2	17.7	72.6	210	2	2	200	239
33	33	11.77	54.1	5.3	17.3	56.0	196	2	1	164	165
34	34	13.59	54.0	6.1	24.2	111.7	312	2	1	258	169
35	35	9.74	54.4	6.3	11.4	76.1	221	2	2	170	172
36	36	10.33	55.8	5.0	21.2	104.3	266	2	1	181	149
37	37	9.97	58.2	2.8	16.5	76.5	90	2	2	69	42
38	38	7.84	49.1	4.6	7.1	87.9	60	2	3	50	45
39	39	10.47	53.2	4.1	5.7	69.1	196	2	2	168	153
40	40	8.16	60.9	1.3	1.9	58.0	73	2	3	49	21
41	41	8.48	51.1	3.7	12.1	92.8	166	2	3	145	118
42	42	10.72	53.8	4.7	23.2	94.1	113	2	3	90	107
43	43	11.20	45.0	3.0	7.0	78.9	130	2	3	95	56
44	44	10.12	51.7	5.6	14.9	79.1	362	1	3	313	264
45	45	8.37	50.7	5.5	15.1	84.8	115	2	2	96	88
46	46	10.16	54.2	4.6	8.4	51.5	831	1	4	581	629
48	48	10.90	57.2	5.5	10.6	71.9	593	2	2	446	211
49	49	7.67	51.7	1.8	2.5	40.4	106	2	3	93	35
50	50	8.88	51.5	4.2	10.1	86.9	305	2	3	238	197
51	51	11.48	57.6	5.6	20.3	82.0	252	2	1	207	251
52	52	9.23	51.6	4.3	11.6	42.6	620	2	2	413	420
53	53	11.41	61.1	7.6	16.6	97.9	535	2	3	330	273
54	54	12.07	43.7	7.8	52.4	105.3	157	2	2	115	76
55	55	8.63	54.0	3.1	8.4	56.2	76	2	1	39	44
56	56	11.15	56.5	3.9	7.7	73.9	281	2	1	217	199
57	57	7.14	59.0	3.7	2.6	75.8	70	2	4	37	35
58	58	7.65	47.1	4.3	16.4	65.7	318	2	4	265	314
59	59	10.73	50.6	3.9	19.3	101.0	445	1	2	374	345
60	60	11.46	56.9	4.5	15.6	97.7	191	2	3	153	132
61	61	10.42	58.0	3.4	8.0	59.0	119	2	1	67	64
62	62	11.18	51.0	5.7	18.8	55.9	595	1	2	546	392
63	63	7.93	64.1	5.4	7.5	98.1	68	2	4	42	49

64	64	9.66	52.1	4.4	9.9	98.3	83	2	2	66	95
65	65	7.78	45.5	5.0	20.9	71.6	489	2	3	391	329
66	66	9.42	50.6	4.3	24.8	62.8	508	2	1	421	528
67	67	10.02	49.5	4.4	8.3	93.0	265	2	2	191	202
68	68	8.58	55.0	3.7	7.4	95.9	304	2	3	248	218
69	69	9.61	52.4	4.5	6.9	87.2	487	2	3	404	220
70	70	8.03	54.2	3.5	24.3	87.3	97	2	1	65	55
71	71	7.39	51.0	4.2	14.6	88.4	72	2	2	38	67
72	72	7.08	52.0	2.0	12.3	56.4	87	2	3	52	57
73	73	9.53	51.5	5.2	15.0	65.7	298	2	3	241	193
74	74	10.05	52.0	4.5	36.7	87.5	184	1	1	144	151
75	75	8.45	38.8	3.4	12.9	85.0	235	2	2	143	124
76	76	6.70	48.6	4.5	13.0	80.8	76	2	4	51	79
77	77	8.90	49.7	2.9	12.7	86.9	52	2	1	37	35
78	78	10.23	53.2	4.9	9.9	77.9	752	1	2	595	446
79	79	8.88	55.8	4.4	14.1	76.8	237	2	2	165	182
80	80	10.30	59.6	5.1	27.8	88.9	175	2	2	113	73
81	81	10.79	44.2	2.9	2.6	56.6	461	1	2	320	196
82	82	7.94	49.5	3.5	6.2	92.3	195	2	2	139	116
83	83	7.63	52.1	5.5	11.6	61.1	197	2	4	109	110
84	84	8.77	54.5	4.7	5.2	47.0	143	2	4	85	87
85	85	8.09	56.9	1.7	7.6	56.9	92	2	3	61	61
86	86	9.05	51.2	4.1	20.5	79.8	195	2	3	127	112
87	87	7.91	52.8	2.9	11.9	79.5	477	2	3	349	188
88	88	10.39	54.6	4.3	14.0	88.3	353	2	2	223	200
89	89	9.36	54.1	4.8	18.3	90.6	165	2	1	127	158
90	90	11.41	50.4	5.8	23.8	73.0	424	1	3	359	335
91	91	8.86	51.3	2.9	9.5	87.5	100	2	3	65	53
92	92	8.93	56.0	2.0	6.2	72.5	95	2	3	59	56
93	93	8.92	53.9	1.3	2.2	79.5	56	2	2	40	14
94	94	8.15	54.9	5.3	12.3	79.8	99	2	4	55	71
95	95	9.77	50.2	5.3	15.7	89.7	154	2	2	123	148
96	96	8.54	56.1	2.5	27.0	82.5	98	2	1	57	75
97	97	8.66	52.8	3.8	6.8	69.5	246	2	3	178	177
98	98	12.01	52.8	4.8	10.8	96.9	298	2	1	237	115
99	99	7.95	51.8	2.3	4.6	54.9	163	2	3	128	93
100	100	10.15	51.9	6.2	16.4	59.2	568	1	3	452	371
101	101	9.76	53.2	2.6	6.9	80.1	64	2	4	47	55
102	102	9.89	45.2	4.3	11.8	108.7	190	2	1	141	112
103	103	7.14	57.6	2.7	13.1	92.6	92	2	4	40	50
104	104	13.95	65.9	6.6	15.6	133.5	356	2	1	308	182
105	105	9.44	52.5	4.5	10.9	58.5	297	2	3	230	263
106	106	10.80	63.9	2.9	1.6	57.4	130	2	3	69	62

107	107	7.14	51.7	1.4	4.1	45.7	115	2	3	90	19
108	108	8.02	55.0	2.1	3.8	46.5	91	2	2	44	32
109	109	11.80	53.8	5.7	9.1	116.9	571	1	2	441	469
110	110	9.50	49.3	5.8	42.0	70.9	98	2	3	68	46
111	111	7.70	56.9	4.4	12.2	67.9	129	2	4	85	136
113	113	9.41	59.5	3.1	20.6	91.7	29	2	3	20	22

Facilities

1	60.0
2	40.0
3	20.0
4	40.0
5	40.0
6	40.0
7	40.0
8	60.0
9	40.0
10	40.0
11	80.0
12	40.0
13	57.1
14	37.1
15	17.1
16	57.1
17	37.1
18	37.1
19	37.1
20	77.1
21	17.1
22	37.1
23	57.1
24	57.1
25	57.1
26	57.1
27	37.1
28	37.1
29	37.1
30	37.1
31	57.1
32	54.3
33	34.3
34	54.3
35	54.3
36	54.3

37	34.3
38	34.3
39	54.3
40	14.3
41	34.3
42	34.3
43	34.3
44	54.3
45	34.3
46	74.3
48	51.4
49	11.4
50	51.4
51	51.4
52	71.4
53	51.4
54	31.4
55	31.4
56	51.4
57	31.4
58	51.4
59	51.4
60	31.4
61	31.4
62	68.6
63	28.6
64	28.6
65	48.6
66	48.6
67	48.6
68	48.6
69	48.6
70	28.6
71	28.6
72	28.6
73	48.6
74	68.6
75	48.6
76	28.6
77	28.6
78	68.6
79	48.6
80	45.7

81	65.7
82	45.7
83	45.7
84	25.7
85	45.7
86	45.7
87	65.7
88	65.7
89	45.7
90	45.7
91	25.7
92	25.7
93	5.7
94	25.7
95	25.7
96	45.7
97	45.7
98	45.7
99	42.9
100	62.9
101	22.9
102	42.9
103	22.9
104	62.9
105	42.9
106	22.9
107	22.9
108	22.9
109	62.9
110	22.9
111	62.9
113	22.9

20. Create new dummy/indicator columns (i1, i2, i3, i4) for regions using `ifelse()` function. For example, `i1 = 1` when `Region = 1` and `i1 = 0` when `Region` is not equal to 1; `i2 = 1` when `Region = 2` and `i2 = 0` when `Region` is not equal to 2; ...

```
infectionrisk$i1 <- ifelse(infectionrisk$Region == 1, 1, 0)
infectionrisk$i2 <- ifelse(infectionrisk$Region == 2, 1, 0)
infectionrisk$i3 <- ifelse(infectionrisk$Region == 3, 1, 0)
infectionrisk$i4 <- ifelse(infectionrisk$Region == 4, 1, 0)
```

21. Fit a multiple linear regression model of `InfctRsk` on `Stay + Xray + i2 + i3 + i4`.

```
model.7 <- lm(InfctRsk ~ Stay + Xray + i2 + i3 + i4, data=infectionrisk)
summary(model.7)
```

Call:

```
lm(formula = InfctRsk ~ Stay + Xray + i2 + i3 + i4, data = infectionrisk)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-2.66492	-0.65420	0.04265	0.64034	2.51391

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-2.134259	0.877347	-2.433	0.01668	*
Stay	0.505394	0.081455	6.205	1.11e-08	***
Xray	0.017587	0.005649	3.113	0.00238	**
i2	0.171284	0.281475	0.609	0.54416	
i3	0.095461	0.288852	0.330	0.74169	
i4	1.057835	0.378077	2.798	0.00612	**

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.036 on 105 degrees of freedom

Multiple R-squared: 0.4198, Adjusted R-squared: 0.3922

F-statistic: 15.19 on 5 and 105 DF, p-value: 3.243e-11

23. Can we include i1 + i2 + i3 + i4 in this multiple linear regression? Why?

No. In the context of using dummy variables for categorical data in regression analysis, it's essential to designate a reference category. This reference category is represented by a coefficient of zero, while the coefficients for the other categories are interpreted as deviations from this reference point.

```
model.8 <- lm(InfctRsk ~ Stay + Xray + i1 + i2 + i3 + i4, data=infectionrisk)
summary(model.8)
```

Call:

```
lm(formula = InfctRsk ~ Stay + Xray + i1 + i2 + i3 + i4, data = infectionrisk)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-2.66492	-0.65420	0.04265	0.64034	2.51391

Coefficients: (1 not defined because of singularities)

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-1.076424	0.721361	-1.492	0.13864
Stay	0.505394	0.081455	6.205	1.11e-08 ***
Xray	0.017587	0.005649	3.113	0.00238 **
i1	-1.057835	0.378077	-2.798	0.00612 **
i2	-0.886551	0.339887	-2.608	0.01042 *
i3	-0.962374	0.323365	-2.976	0.00362 **
i4	NA	NA	NA	NA

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.036 on 105 degrees of freedom

Multiple R-squared: 0.4198, Adjusted R-squared: 0.3922

F-statistic: 15.19 on 5 and 105 DF, p-value: 3.243e-11

24. Conduct an F-test (use `anova()` function) to see if at least one of i2, i3, and i4 are useful.

```
model.9 <- lm(InfctRsk ~ Stay + Xray, data=infectionrisk)
anova(model.7, model.9)
```

Analysis of Variance Table

Model 1: InfctRsk ~ Stay + Xray + i2 + i3 + i4

Model 2: InfctRsk ~ Stay + Xray

	Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
1	105	112.71				
2	108	123.56	-3	-10.849	3.3687	0.02135 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Logistic regression

Simple example

```
logit_m2 =glm(formula = LUNG_CANCER ~ ANXIETY+PEER_PRESSURE+`CHRONIC DISEASE`+FATIGUE+ALLERGY
```

Confusion Matrix

Total population = P + N	Predicted Positive (PP)	Predicted Negative (PN)	Informedness, bookmaker informedness (BM) = TPR + TNR - 1	Prevalence threshold (PT) = $\frac{\sqrt{TPR \times FPR} - FPR}{TPR - FPR}$
Positive (P) ^[a]	True positive (TP), hit ^[b]	False negative (FN), miss, underestimation	True positive rate (TPR), recall, sensitivity (SEN), probability of detection, hit rate, power = $\frac{TP}{P} = 1 - FNR$	False negative rate (FNR), miss rate type II error ^[c] = $\frac{FN}{P} = 1 - TPR$
Negative (N) ^[d]	False positive (FP), false alarm, overestimation	True negative (TN), correct rejection ^[e]	False positive rate (FPR), probability of false alarm, fall-out type I error ^[7] = $\frac{FP}{N} = 1 - TNR$	True negative rate (TNR), specificity (SPC), selectivity = $\frac{TN}{N} = 1 - FPR$

example

Lung Cancer Classification (<https://www.kaggle.com/datasets/mysarahmadbhat/lung-cancer?select=survey+lung+cancer.csv>)

The effectiveness of cancer prediction system helps the people to know their cancer risk with low cost and it also helps the people to take the appropriate decision based on their cancer risk status. The data is collected from the website online lung cancer prediction system.

Total no. of attributes: 16 No. of instances: 284

Attribute information:

Gender: M(male), F(female) Age: Age of the patient Smoking: YES=2, NO=1. Yellow fingers: YES=2, NO=1. Anxiety: YES=2, NO=1. Peer_pressure: YES=2, NO=1. Chronic Disease: YES=2, NO=1. Fatigue: YES=2, NO=1. Allergy: YES=2, NO=1. Wheezing: YES=2, NO=1. Alcohol: YES=2, NO=1. Coughing: YES=2, NO=1. Shortness of Breath: YES=2, NO=1. Swallowing Difficulty: YES=2, NO=1. Chest pain: YES=2, NO=1. Lung Cancer: YES, NO.

Goal: It is your job to classify Lung Cancer using other variables

Example Code

```
#Load the dataset
library(readr)
data = read_csv('/Users/luyu/Desktop/survey_lung_cancer.csv', show_col_types = FALSE)
data$LUNG_CANCER <- ifelse(data$LUNG_CANCER=="YES", 1, 0)
summary(data)
```

GENDER		AGE		SMOKING		YELLOW_FINGERS	
Length:	309	Min.	:21.00	Min.	:1.000	Min.	:1.00
Class:	character	1st Qu.:	:57.00	1st Qu.:	:1.000	1st Qu.:	:1.00
Mode	:character	Median	:62.00	Median	:2.000	Median	:2.00
		Mean	:62.67	Mean	:1.563	Mean	:1.57
		3rd Qu.:	:69.00	3rd Qu.:	:2.000	3rd Qu.:	:2.00
		Max.	:87.00	Max.	:2.000	Max.	:2.00
ANXIETY		PEER_PRESSURE		CHRONIC DISEASE		FATIGUE	
Min.	:1.000	Min.	:1.000	Min.	:1.000	Min.	:1.000
1st Qu.:	:1.000	1st Qu.:	:1.000	1st Qu.:	:1.000	1st Qu.:	:1.000
Median	:1.000	Median	:2.000	Median	:2.000	Median	:2.000
Mean	:1.498	Mean	:1.502	Mean	:1.505	Mean	:1.673
3rd Qu.:	:2.000	3rd Qu.:	:2.000	3rd Qu.:	:2.000	3rd Qu.:	:2.000
Max.	:2.000	Max.	:2.000	Max.	:2.000	Max.	:2.000
ALLERGY		WHEEZING		ALCOHOL CONSUMING		COUGHING	
Min.	:1.000	Min.	:1.000	Min.	:1.000	Min.	:1.000
1st Qu.:	:1.000	1st Qu.:	:1.000	1st Qu.:	:1.000	1st Qu.:	:1.000
Median	:2.000	Median	:2.000	Median	:2.000	Median	:2.000
Mean	:1.557	Mean	:1.557	Mean	:1.557	Mean	:1.579
3rd Qu.:	:2.000	3rd Qu.:	:2.000	3rd Qu.:	:2.000	3rd Qu.:	:2.000
Max.	:2.000	Max.	:2.000	Max.	:2.000	Max.	:2.000
SHORTNESS OF BREATH		SWALLOWING DIFFICULTY		CHEST PAIN		LUNG_CANCER	
Min.	:1.000	Min.	:1.000	Min.	:1.000	Min.	:0.0000
1st Qu.:	:1.000	1st Qu.:	:1.000	1st Qu.:	:1.000	1st Qu.:	:1.0000
Median	:2.000	Median	:1.000	Median	:2.000	Median	:1.0000
Mean	:1.641	Mean	:1.469	Mean	:1.557	Mean	:0.8738
3rd Qu.:	:2.000	3rd Qu.:	:2.000	3rd Qu.:	:2.000	3rd Qu.:	:1.0000
Max.	:2.000	Max.	:2.000	Max.	:2.000	Max.	:1.0000

```
library(ggplot2)
ggplot(data, aes(x = factor(SMOKING), fill = factor(LUNG_CANCER))) +
  geom_bar(position = "fill") +
  scale_fill_manual(values = c("0" = "lightblue", "1" = "salmon")) +
  theme_minimal()
```

```
### Data SAMPLING ###
library(caret)
```

Attaching package: 'caret'

The following object is masked from 'package:purrr':

lift

The following object is masked from 'package:mosaic':

dotPlot

The following object is masked from 'package:survival':

cluster

```
set.seed(101)
split = createDataPartition(data$LUNG_CANCER, p = 0.80, list = FALSE)
train_data = data[split,]
test_data = data[-split,]
nrow(train_data)
```

[1] 248

```
nrow(test_data)
```

[1] 61

```
#error metrics -- Confusion Matrix
err_metric=function(CM)
{
  TN =CM[1,1]
  TP =CM[2,2]
  FP =CM[1,2]
  FN =CM[2,1]
  precision =(TP)/(TP+FP)
  recall_score =(TP)/(TP+FN)
  f1_score=2*((precision*recall_score)/(precision+recall_score))
  accuracy_model =(TP+TN)/(TP+TN+FP+FN)
  False_positive_rate =(FP)/(FP+TN)
  False_negative_rate =(FN)/(FN+TP)
  print(paste("Precision value of the model: ",round(precision,2)))
  print(paste("Accuracy of the model: ",round(accuracy_model,2)))
  print(paste("Recall value of the model: ",round(recall_score,2)))
  print(paste("False Positive rate of the model: ",round(False_positive_rate,2)))
  print(paste("False Negative rate of the model: ",round(False_negative_rate,2)))
  print(paste("F1 score of the model: ",round(f1_score,2)))
}
```

```
# Logistic regression
logit_m = glm(formula = LUNG_CANCER ~ ., data = train_data, family = 'binomial')
summary(logit_m)
```

Call:

```
glm(formula = LUNG_CANCER ~ ., family = "binomial", data = train_data)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-32.13449	6.57393	-4.888	1.02e-06 ***
GENDERM	-0.84716	0.82793	-1.023	0.306203
AGE	0.01537	0.03490	0.440	0.659756
SMOKING	1.10495	0.82246	1.343	0.179119
YELLOW_FINGERS	1.10971	0.82005	1.353	0.175982
ANXIETY	2.08645	1.08610	1.921	0.054725 .
PEER_PRESSURE	1.94159	0.74009	2.623	0.008704 **
`CHRONIC DISEASE`	3.91378	1.12190	3.489	0.000486 ***
FATIGUE	2.62906	0.89032	2.953	0.003148 **
ALLERGY	1.42702	0.83764	1.704	0.088454 .
WHEEZING	1.07933	0.90761	1.189	0.234357
`ALCOHOL CONSUMING`	2.41116	0.98495	2.448	0.014365 *
COUGHING	3.14783	1.22386	2.572	0.010110 *
`SHORTNESS OF BREATH`	-0.22681	0.84245	-0.269	0.787757
`SWALLOWING DIFFICULTY`	2.24613	1.24347	1.806	0.070865 .
`CHEST PAIN`	0.89356	0.73930	1.209	0.226791

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 198.230 on 247 degrees of freedom
Residual deviance: 77.353 on 232 degrees of freedom
AIC: 109.35

Number of Fisher Scoring iterations: 8

```
# Logistic regression
logit_m2 = glm(formula = LUNG_CANCER ~ ANXIETY+PEER_PRESSURE+`CHRONIC DISEASE`+FATIGUE+ALLERGY+`ALCOHOL CONSUMING`+COUGHING+`SWALLOWING DIFFICULTY`,
summary(logit_m2)
```

Call:

```
glm(formula = LUNG_CANCER ~ ANXIETY + PEER_PRESSURE + `CHRONIC DISEASE` +
    FATIGUE + ALLERGY + `ALCOHOL CONSUMING` + COUGHING + `SWALLOWING DIFFICULTY`,
    family = "binomial", data = train_data)
```

```

Coefficients:
                Estimate Std. Error z value Pr(>|z|)
(Intercept)      -27.3830    5.5206  -4.960 7.05e-07 ***
ANXIETY            2.5514    0.8659   2.947 0.003213 **
PEER_PRESSURE      2.1822    0.7245   3.012 0.002593 **
`CHRONIC DISEASE`  3.5120    0.9696   3.622 0.000292 ***
FATIGUE            2.4939    0.6979   3.573 0.000353 ***
ALLERGY            2.0104    0.7428   2.707 0.006796 **
`ALCOHOL CONSUMING` 2.5084    0.8653   2.899 0.003744 **
COUGHING           3.2220    0.9540   3.377 0.000732 ***
`SWALLOWING DIFFICULTY` 2.5273    1.0096   2.503 0.012309 *

```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 198.230 on 247 degrees of freedom
Residual deviance: 83.501 on 239 degrees of freedom
AIC: 101.5

Number of Fisher Scoring iterations: 8

```

library(dplyr)
logit_P_prob = predict(logit_m, newdata = select(test_data, -LUNG_CANCER), type = 'response')
logit_P_prob[1:3]

```

```

      1      2      3
0.9912600 0.9997708 0.9980692

```

```

logit_P <- ifelse(logit_P_prob > 0.5, 1, 0) # Probability check
logit_P[1:3]

```

```

1 2 3
1 1 1

```

```

CM = table(test_data$LUNG_CANCER, logit_P)
print(CM)

```

```

logit_P
  0  1
0  3  2
1  1 55

```



```
err_metric(CM)
```

```
[1] "Precision value of the model: 0.96"  
[1] "Accuracy of the model: 0.95"  
[1] "Recall value of the model: 0.98"  
[1] "False Positive rate of the model: 0.4"  
[1] "False Negative rate of the model: 0.02"  
[1] "F1 score of the model: 0.97"
```

```
#ROC-curve using pROC library  
library(pROC)
```

Type 'citation("pROC")' for a citation.

Attaching package: 'pROC'

The following objects are masked from 'package:mosaic':

cov, var

The following objects are masked from 'package:stats':

cov, smooth, var

```
roc_score=roc(test_data$LUNG_CANCER, logit_P_prob) #AUC score
```

Setting levels: control = 0, case = 1

Setting direction: controls < cases

```
plot(roc_score, main = "ROC curve -- Logistic Regression")
```

